

In [1]:

```
from keras.utils import np_utils
from keras.datasets import mnist
import seaborn as sns
from keras.initializers import RandomNormal
```

Using TensorFlow backend.

In [2]:

```
#!/usr/bin/env python
%matplotlib notebook
%matplotlib inline
%pylab inline
import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    fig.canvas.draw()
```

Populating the interactive namespace from numpy and matplotlib

## IMPORTING AND DIVING INTO THE DATASET

In [3]:

```
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

Downloading data from <https://s3.amazonaws.com/img-datasets/mnist.npz>  
11493376/11490434 [=====] - 1s 0us/step

In [4]:

```
print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d, %d)"%(X_train.shape[1], X_train.shape[2]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d, %d)"%(X_test.shape[1], X_test.shape[2]))
```


Number of training examples : 60000 and each image is of shape (28, 28)


Number of training examples : 10000 and each image is of shape (28, 28)

In [5]:

```
plt.imshow(X_train[2])
print("The output is", y_train[2])
```

The output is 4





In [0]:

```
X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])
```

In [7]:

```
print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d)
"%(X_train.shape[1]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d)"
"%(X_test.shape[1]))
```

Number of training examples : 60000 and each image is of shape (784)  
Number of training examples : 10000 and each image is of shape (784)

In [0]:

```
X_train = X_train/255
X_test = X_test/255
```

In [9]:

```
# here we are having a class number for each image
print("Class label of first image :", y_train[0])

# lets convert this into a 10 dimensional vector
# ex: consider an image is 5 convert it into 5 => [0, 0, 0, 0, 0, 1, 0, 0, 0, 0]
# this conversion needed for MLPs

Y_train = np_utils.to_categorical(y_train, 10)
Y_test = np_utils.to_categorical(y_test, 10)

print("After converting the output into a vector : ",Y_train[0])
```

Class label of first image : 5  
After converting the output into a vector : [0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]

In [0]:

```
from keras.models import Sequential
from keras.layers import Dense, Activation
```

In [0]:

```
output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128
nb_epoch = 20
```

## MODEL 1 - 2 HIDDEN LAYERS

### 1.1 RELU ACTIVATION

In [0]:

```
model1 = Sequential()
model1.add(Dense(230, input_dim=input_dim, activation='relu'))
model1.add(Dense(100, input_dim=input_dim, activation='relu'))
model1.add(Dense(10, input_dim=input_dim, activation='softmax'))
```

WARNING: Logging before flag parsing goes to stderr.  
W0822 15:13:09.473532 140115941365632 deprecation\_wrapper.py:119] From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:74: The name tf.get\_default\_graph is deprecated. Please use tf.compat.v1.get\_default\_graph instead.

```
W0822 15:13:09.522732 140115941365632 deprecation_wrapper.py:119] From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:517: The name tf.placeholder is deprecated. Please use tf.compat.v1.placeholder instead.
```

```
W0822 15:13:09.530593 140115941365632 deprecation_wrapper.py:119] From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:4138: The name tf.random_uniform is deprecated. Please use tf.random.uniform instead.
```

In [0]:

```
model1.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model1.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1,
                    validation_data=(X_test, Y_test))
```

```
W0822 15:13:14.105819 140115941365632 deprecation_wrapper.py:119] From /usr/local/lib/python3.6/dist-packages/keras/optimizers.py:790: The name tf.train.Optimizer is deprecated. Please use tf.compat.v1.train.Optimizer instead.
```

```
W0822 15:13:14.151415 140115941365632 deprecation_wrapper.py:119] From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:3295: The name tf.log is deprecated. Please use tf.math.log instead.
```

```
W0822 15:13:14.289976 140115941365632 deprecation.py:323] From /usr/local/lib/python3.6/dist-packages/tensorflow/python/ops/math_grad.py:1250: add_dispatch_support.<locals>.wrapper (from tensorflow.python.ops.array_ops) is deprecated and will be removed in a future version.
```

Instructions for updating:

Use tf.where in 2.0, which has the same broadcast rule as np.where

```
W0822 15:13:14.354537 140115941365632 deprecation_wrapper.py:119] From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:986: The name tf.assign_add is deprecated. Please use tf.compat.v1.assign_add instead.
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

```
60000/60000 [=====] - 7s 116us/step - loss: 0.2841 - acc: 0.9194
- val_loss: 0.1420 - val_acc: 0.9564
```

Epoch 2/20

```
60000/60000 [=====] - 3s 48us/step - loss: 0.1106 - acc: 0.9673
- val_loss: 0.0923 - val_acc: 0.9710
```

Epoch 3/20

```
60000/60000 [=====] - 3s 49us/step - loss: 0.0713 - acc: 0.9783
- val_loss: 0.0827 - val_acc: 0.9735
```

Epoch 4/20

```
60000/60000 [=====] - 3s 49us/step - loss: 0.0525 - acc: 0.9837
- val_loss: 0.0799 - val_acc: 0.9741
```

Epoch 5/20

```
60000/60000 [=====] - 3s 50us/step - loss: 0.0382 - acc: 0.9885
- val_loss: 0.0801 - val_acc: 0.9764
```

Epoch 6/20

```
60000/60000 [=====] - 3s 45us/step - loss: 0.0286 - acc: 0.9913
- val_loss: 0.0762 - val_acc: 0.9771
```

Epoch 7/20

```
60000/60000 [=====] - 3s 45us/step - loss: 0.0231 - acc: 0.9926
- val_loss: 0.0733 - val_acc: 0.9780
```

Epoch 8/20

```
60000/60000 [=====] - 3s 45us/step - loss: 0.0202 - acc: 0.9936
- val_loss: 0.0704 - val_acc: 0.9802
```

Epoch 9/20

```
60000/60000 [=====] - 3s 47us/step - loss: 0.0160 - acc: 0.9949
- val_loss: 0.0788 - val_acc: 0.9785
```

Epoch 10/20

```
60000/60000 [=====] - 3s 45us/step - loss: 0.0124 - acc: 0.9961
- val_loss: 0.0883 - val_acc: 0.9782
```

Epoch 11/20

```
60000/60000 [=====] - 3s 47us/step - loss: 0.0141 - acc: 0.9952
- val_loss: 0.0886 - val_acc: 0.9782
```

Epoch 12/20

```
60000/60000 [=====] - 3s 46us/step - loss: 0.0132 - acc: 0.9957
- val_loss: 0.0937 - val_acc: 0.9764
```

```

val_loss: 0.0780 - val_acc: 0.9815
Epoch 13/20
60000/60000 [=====] - 3s 47us/step - loss: 0.0089 - acc: 0.9971
- val_loss: 0.0780 - val_acc: 0.9815
Epoch 14/20
60000/60000 [=====] - 3s 45us/step - loss: 0.0071 - acc: 0.9975
- val_loss: 0.0841 - val_acc: 0.9802
Epoch 15/20
60000/60000 [=====] - 3s 46us/step - loss: 0.0137 - acc: 0.9955
- val_loss: 0.0879 - val_acc: 0.9793
Epoch 16/20
60000/60000 [=====] - 3s 45us/step - loss: 0.0084 - acc: 0.9971
- val_loss: 0.0865 - val_acc: 0.9788
Epoch 17/20
60000/60000 [=====] - 3s 46us/step - loss: 0.0058 - acc: 0.9982
- val_loss: 0.0994 - val_acc: 0.9786
Epoch 18/20
60000/60000 [=====] - 3s 45us/step - loss: 0.0116 - acc: 0.9962
- val_loss: 0.0798 - val_acc: 0.9807
Epoch 19/20
60000/60000 [=====] - 3s 45us/step - loss: 0.0072 - acc: 0.9975
- val_loss: 0.0874 - val_acc: 0.9802
Epoch 20/20
60000/60000 [=====] - 3s 45us/step - loss: 0.0047 - acc: 0.9986
- val_loss: 0.0807 - val_acc: 0.9809

```

In [0]:

```
model1.summary()
```

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 230)	180550
dense_2 (Dense)	(None, 100)	23100
dense_3 (Dense)	(None, 10)	1010

Total params: 204,660  
 Trainable params: 204,660  
 Non-trainable params: 0

In [0]:

```

score = model1.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs


vy = history.history['val_loss']
ty = history.history['loss']

```

```
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.08070132349427932

Test accuracy: 0.9809



The plot says that Validation loss is very high than training loss after 3 epochs

In [0]:


```
w_after = model1.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure(figsize=(10,7))
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```





## 1.2 MODEL WITH 2 POWERS AS NUMBER OF HIDDEN LAYERS (RELU ACTIVATION)

In [0]:

```
model1 = Sequential()
model1.add(Dense(256, input_dim=input_dim, activation='relu'))
model1.add(Dense(128, input_dim=input_dim, activation='relu'))
model1.add(Dense(10, input_dim=input_dim, activation='softmax'))
```

In [0]:

```
model1.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model1.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1,
                    validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 3s 51us/step - loss: 0.2706 - acc: 0.9233
- val_loss: 0.1266 - val_acc: 0.9618
Epoch 2/20
60000/60000 [=====] - 3s 45us/step - loss: 0.1033 - acc: 0.9690
- val_loss: 0.0958 - val_acc: 0.9689
Epoch 3/20
60000/60000 [=====] - 3s 47us/step - loss: 0.0681 - acc: 0.9796
- val_loss: 0.0765 - val_acc: 0.9755
Epoch 4/20
60000/60000 [=====] - 3s 45us/step - loss: 0.0479 - acc: 0.9853
- val_loss: 0.0738 - val_acc: 0.9760
Epoch 5/20
60000/60000 [=====] - 3s 47us/step - loss: 0.0362 - acc: 0.9885
- val_loss: 0.0722 - val_acc: 0.9778
Epoch 6/20
60000/60000 [=====] - 3s 45us/step - loss: 0.0280 - acc: 0.9912
- val_loss: 0.0737 - val_acc: 0.9783
Epoch 7/20
60000/60000 [=====] - 3s 46us/step - loss: 0.0204 - acc: 0.9937
- val_loss: 0.0681 - val_acc: 0.9804
Epoch 8/20
60000/60000 [=====] - 3s 44us/step - loss: 0.0195 - acc: 0.9934
- val_loss: 0.0733 - val_acc: 0.9795
Epoch 9/20
60000/60000 [=====] - 3s 46us/step - loss: 0.0141 - acc: 0.9953
- val_loss: 0.0857 - val_acc: 0.9777
Epoch 10/20
60000/60000 [=====] - 3s 45us/step - loss: 0.0148 - acc: 0.9953
- val_loss: 0.0783 - val_acc: 0.9786
Epoch 11/20
60000/60000 [=====] - 3s 47us/step - loss: 0.0100 - acc: 0.9969
- val_loss: 0.1074 - val_acc: 0.9727
Epoch 12/20
60000/60000 [=====] - 3s 47us/step - loss: 0.0123 - acc: 0.9957
- val_loss: 0.0991 - val_acc: 0.9750
Epoch 13/20
60000/60000 [=====] - 3s 47us/step - loss: 0.0098 - acc: 0.9967
- val_loss: 0.0883 - val_acc: 0.9784
Epoch 14/20
60000/60000 [=====] - 3s 47us/step - loss: 0.0077 - acc: 0.9975
- val_loss: 0.0950 - val_acc: 0.9782
Epoch 15/20
60000/60000 [=====] - 3s 46us/step - loss: 0.0098 - acc: 0.9965
- val_loss: 0.1035 - val_acc: 0.9773
Epoch 16/20
60000/60000 [=====] - 3s 46us/step - loss: 0.0059 - acc: 0.9980
- val_loss: 0.0860 - val_acc: 0.9803
Epoch 17/20
```

```
Epoch 17/20
60000/60000 [=====] - 3s 46us/step - loss: 0.0090 - acc: 0.9966
- val_loss: 0.0943 - val_acc: 0.9784
Epoch 18/20
60000/60000 [=====] - 3s 46us/step - loss: 0.0092 - acc: 0.9970
- val_loss: 0.1059 - val_acc: 0.9787
Epoch 19/20
60000/60000 [=====] - 3s 47us/step - loss: 0.0056 - acc: 0.9980
- val_loss: 0.1003 - val_acc: 0.9781
Epoch 20/20
60000/60000 [=====] - 3s 46us/step - loss: 0.0067 - acc: 0.9978
- val_loss: 0.0951 - val_acc: 0.9802
```

In [0]:

```
model1.summary()
```

Layer (type)	Output Shape	Param #
dense_4 (Dense)	(None, 256)	200960
dense_5 (Dense)	(None, 128)	32896
dense_6 (Dense)	(None, 10)	1290
Total params: 235,146		
Trainable params: 235,146		
Non-trainable params: 0		

In [0]:

```
score = model1.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))


# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.09506374017388317  
Test accuracy: 0.9802





**The model overfits as the number of epochs increased.**

In [0]:


```
w_after = model1.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure(figsize=(10,7))
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



### 1.3 RELU ACTIVATION WITH DROPOUT VALUE 0.5

In [0]:

```
from keras.layers import Dropout
model1 = Sequential()
```



```

model1.add(Dense(256, input_dim=input_dim, activation='relu'))
model1.add(Dropout(0.5))
model1.add(Dense(128, input_dim=input_dim, activation='relu'))
model1.add(Dropout(0.5))
model1.add(Dense(10, input_dim=input_dim, activation='softmax'))

```

W0822 15:32:48.802173 140115941365632 deprecation.py:506] From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:3445: calling dropout (from tensorflow.python.ops.nn\_ops) with keep\_prob is deprecated and will be removed in a future version. Instructions for updating:  
Please use `rate` instead of `keep\_prob`. Rate should be set to `rate = 1 - keep\_prob`.

In [0]:

```

model1.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model1.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1,
                    validation_data=(X_test, Y_test))

```

Train on 60000 samples, validate on 10000 samples

```

Epoch 1/20
60000/60000 [=====] - 3s 53us/step - loss: 0.5257 - acc: 0.8390
- val_loss: 0.1702 - val_acc: 0.9497
Epoch 2/20
60000/60000 [=====] - 3s 49us/step - loss: 0.2422 - acc: 0.9304
- val_loss: 0.1223 - val_acc: 0.9646
Epoch 3/20
60000/60000 [=====] - 3s 49us/step - loss: 0.1883 - acc: 0.9457
- val_loss: 0.1009 - val_acc: 0.9699
Epoch 4/20
60000/60000 [=====] - 3s 49us/step - loss: 0.1653 - acc: 0.9517
- val_loss: 0.0946 - val_acc: 0.9704
Epoch 5/20
60000/60000 [=====] - 3s 48us/step - loss: 0.1470 - acc: 0.9568
- val_loss: 0.0860 - val_acc: 0.9743
Epoch 6/20
60000/60000 [=====] - 3s 49us/step - loss: 0.1327 - acc: 0.9614
- val_loss: 0.0816 - val_acc: 0.9741
Epoch 7/20
60000/60000 [=====] - 3s 48us/step - loss: 0.1190 - acc: 0.9655
- val_loss: 0.0778 - val_acc: 0.9763
Epoch 8/20
60000/60000 [=====] - 3s 50us/step - loss: 0.1144 - acc: 0.9663
- val_loss: 0.0796 - val_acc: 0.9767
Epoch 9/20
60000/60000 [=====] - 3s 48us/step - loss: 0.1060 - acc: 0.9682
- val_loss: 0.0748 - val_acc: 0.9791
Epoch 10/20
60000/60000 [=====] - 3s 49us/step - loss: 0.1007 - acc: 0.9694
- val_loss: 0.0720 - val_acc: 0.9782
Epoch 11/20
60000/60000 [=====] - 3s 49us/step - loss: 0.0958 - acc: 0.9709
- val_loss: 0.0708 - val_acc: 0.9783
Epoch 12/20
60000/60000 [=====] - 3s 48us/step - loss: 0.0920 - acc: 0.9721
- val_loss: 0.0711 - val_acc: 0.9784
Epoch 13/20
60000/60000 [=====] - 3s 49us/step - loss: 0.0896 - acc: 0.9721
- val_loss: 0.0692 - val_acc: 0.9798
Epoch 14/20
60000/60000 [=====] - 3s 48us/step - loss: 0.0870 - acc: 0.9730
- val_loss: 0.0687 - val_acc: 0.9798
Epoch 15/20
60000/60000 [=====] - 3s 48us/step - loss: 0.0833 - acc: 0.9742
- val_loss: 0.0656 - val_acc: 0.9809
Epoch 16/20
60000/60000 [=====] - 3s 48us/step - loss: 0.0781 - acc: 0.9761
- val_loss: 0.0676 - val_acc: 0.9811
Epoch 17/20
60000/60000 [=====] - 3s 49us/step - loss: 0.0772 - acc: 0.9762
- val_loss: 0.0688 - val_acc: 0.9808
Epoch 18/20

```

```
60000/60000 [=====] - 3s 49us/step - loss: 0.0745 - acc: 0.9771
- val_loss: 0.0700 - val_acc: 0.9803
Epoch 19/20
60000/60000 [=====] - 3s 48us/step - loss: 0.0727 - acc: 0.9773
- val_loss: 0.0688 - val_acc: 0.9811
Epoch 20/20
60000/60000 [=====] - 3s 50us/step - loss: 0.0698 - acc: 0.9785
- val_loss: 0.0651 - val_acc: 0.9817
```

In [0]:

```
model1.summary()
```

Layer (type)	Output Shape	Param #
dense_8 (Dense)	(None, 256)	200960
dropout_1 (Dropout)	(None, 256)	0
dense_9 (Dense)	(None, 128)	32896
dropout_2 (Dropout)	(None, 128)	0
dense_10 (Dense)	(None, 10)	1290
Total params: 235,146		
Trainable params: 235,146		
Non-trainable params: 0		

In [0]:

```
score = model1.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.06511110941658844

Test accuracy: 0.9817





This is so far the best model as both validation and train losses met their minimum values at the 20th epoch

In [0]:


```
w_after = model1.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure(figsize=(10,7))
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



## 1.4 RELU ACTIVATION WITH BATCH NORM

In [0]:

```
from keras.layers.normalization import BatchNormalization
model1 = Sequential()
```

```
model1.add(Dense(256, input_dim=input_dim, activation='relu'))
model1.add(BatchNormalization())
model1.add(Dense(128, input_dim=input_dim, activation='relu'))
model1.add(BatchNormalization())
model1.add(Dense(10, input_dim=input_dim, activation='softmax'))
```

In [0]:

```
model1.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model1.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1
, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 5s 90us/step - loss: 0.2043 - acc: 0.9397  
- val\_loss: 0.1063 - val\_acc: 0.9658

Epoch 2/20

60000/60000 [=====] - 5s 78us/step - loss: 0.0804 - acc: 0.9761  
- val\_loss: 0.0876 - val\_acc: 0.9730

Epoch 3/20

60000/60000 [=====] - 5s 79us/step - loss: 0.0528 - acc: 0.9832  
- val\_loss: 0.0848 - val\_acc: 0.9734

Epoch 4/20

60000/60000 [=====] - 5s 80us/step - loss: 0.0382 - acc: 0.9881  
- val\_loss: 0.0808 - val\_acc: 0.9757

Epoch 5/20

60000/60000 [=====] - 5s 81us/step - loss: 0.0290 - acc: 0.9905  
- val\_loss: 0.0792 - val\_acc: 0.9766

Epoch 6/20

60000/60000 [=====] - 5s 77us/step - loss: 0.0252 - acc: 0.9919  
- val\_loss: 0.0919 - val\_acc: 0.9735

Epoch 7/20

60000/60000 [=====] - 5s 79us/step - loss: 0.0204 - acc: 0.9932  
- val\_loss: 0.0706 - val\_acc: 0.9791

Epoch 8/20

60000/60000 [=====] - 5s 81us/step - loss: 0.0170 - acc: 0.9946  
- val\_loss: 0.0836 - val\_acc: 0.9779

Epoch 9/20

60000/60000 [=====] - 5s 77us/step - loss: 0.0165 - acc: 0.9946  
- val\_loss: 0.0769 - val\_acc: 0.9793

Epoch 10/20

60000/60000 [=====] - 5s 80us/step - loss: 0.0141 - acc: 0.9952  
- val\_loss: 0.0873 - val\_acc: 0.9769

Epoch 11/20

60000/60000 [=====] - 5s 79us/step - loss: 0.0126 - acc: 0.9960  
- val\_loss: 0.0741 - val\_acc: 0.9807

Epoch 12/20

60000/60000 [=====] - 5s 77us/step - loss: 0.0122 - acc: 0.9959  
- val\_loss: 0.0835 - val\_acc: 0.9785

Epoch 13/20

60000/60000 [=====] - 5s 77us/step - loss: 0.0105 - acc: 0.9963  
- val\_loss: 0.0873 - val\_acc: 0.9779

Epoch 14/20

60000/60000 [=====] - 5s 77us/step - loss: 0.0112 - acc: 0.9963  
- val\_loss: 0.0938 - val\_acc: 0.9776

Epoch 15/20

60000/60000 [=====] - 5s 79us/step - loss: 0.0100 - acc: 0.9968  
- val\_loss: 0.0840 - val\_acc: 0.9801

Epoch 16/20

60000/60000 [=====] - 5s 81us/step - loss: 0.0093 - acc: 0.9968  
- val\_loss: 0.0808 - val\_acc: 0.9813

Epoch 17/20

60000/60000 [=====] - 5s 76us/step - loss: 0.0099 - acc: 0.9966  
- val\_loss: 0.0896 - val\_acc: 0.9772

Epoch 18/20

60000/60000 [=====] - 5s 76us/step - loss: 0.0090 - acc: 0.9970  
- val\_loss: 0.0860 - val\_acc: 0.9779

Epoch 19/20

60000/60000 [=====] - 5s 77us/step - loss: 0.0077 - acc: 0.9972  
- val\_loss: 0.0824 - val\_acc: 0.9799

Epoch 20/20

60000/60000 [=====] - 5s 79us/step - loss: 0.0072 - acc: 0.9978  
- val\_loss: 0.0869 - val\_acc: 0.9791

In [0]:

```
model1.summary()
```

Layer (type)	Output Shape	Param #
dense_11 (Dense)	(None, 256)	200960
batch_normalization_1 (Batch Normalization)	(None, 256)	1024
dense_12 (Dense)	(None, 128)	32896
batch_normalization_2 (Batch Normalization)	(None, 128)	512
dense_13 (Dense)	(None, 10)	1290

Total params: 236,682  
Trainable params: 235,914  
Non-trainable params: 768

In [0]:

```
score = model1.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.08686929494800279

Test accuracy: 0.9791



2.5 5.0 7.5 10.0 12.5 15.0 17.5 20.0  
epoch

**This model was well trained but failed in validation test due to overfitting.**

In [0]:


```
w_after = model1.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure(figsize=(10,7))
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



## 1.5 RELU ACTIVATION WITH BATCH NORM AND DROPOUT (0.5)

In [0]:

```
from keras.layers.normalization import BatchNormalization
model1 = Sequential()

model1.add(Dense(256, input_dim=input_dim, activation='relu'))
model1.add(BatchNormalization())
model1.add(Dropout(0.5))

model1.add(Dense(128, input_dim=input_dim, activation='relu'))
```

```
model1.add(BatchNormalization())
model1.add(Dropout(0.5))

model1.add(Dense(10, input_dim=input_dim, activation='softmax'))
```

In [0]:

```
model1.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model1.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1
, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 6s 97us/step - loss: 0.4898 - acc: 0.8517
- val_loss: 0.1640 - val_acc: 0.9482
Epoch 2/20
60000/60000 [=====] - 5s 81us/step - loss: 0.2463 - acc: 0.9267
- val_loss: 0.1271 - val_acc: 0.9610
Epoch 3/20
60000/60000 [=====] - 5s 81us/step - loss: 0.1993 - acc: 0.9407
- val_loss: 0.1082 - val_acc: 0.9670
Epoch 4/20
60000/60000 [=====] - 5s 82us/step - loss: 0.1682 - acc: 0.9491
- val_loss: 0.0979 - val_acc: 0.9697
Epoch 5/20
60000/60000 [=====] - 5s 82us/step - loss: 0.1533 - acc: 0.9545
- val_loss: 0.0899 - val_acc: 0.9727
Epoch 6/20
60000/60000 [=====] - 5s 79us/step - loss: 0.1376 - acc: 0.9582
- val_loss: 0.0816 - val_acc: 0.9744
Epoch 7/20
60000/60000 [=====] - 5s 80us/step - loss: 0.1288 - acc: 0.9610
- val_loss: 0.0803 - val_acc: 0.9759
Epoch 8/20
60000/60000 [=====] - 5s 79us/step - loss: 0.1187 - acc: 0.9636
- val_loss: 0.0761 - val_acc: 0.9768
Epoch 9/20
60000/60000 [=====] - 5s 80us/step - loss: 0.1104 - acc: 0.9663
- val_loss: 0.0719 - val_acc: 0.9782
Epoch 10/20
60000/60000 [=====] - 5s 83us/step - loss: 0.1022 - acc: 0.9679
- val_loss: 0.0687 - val_acc: 0.9797
Epoch 11/20
60000/60000 [=====] - 5s 82us/step - loss: 0.1014 - acc: 0.9691
- val_loss: 0.0727 - val_acc: 0.9775
Epoch 12/20
60000/60000 [=====] - 5s 81us/step - loss: 0.0991 - acc: 0.9692
- val_loss: 0.0685 - val_acc: 0.9789
Epoch 13/20
60000/60000 [=====] - 5s 79us/step - loss: 0.0927 - acc: 0.9716
- val_loss: 0.0718 - val_acc: 0.9778
Epoch 14/20
60000/60000 [=====] - 5s 81us/step - loss: 0.0871 - acc: 0.9731
- val_loss: 0.0659 - val_acc: 0.9799
Epoch 15/20
60000/60000 [=====] - 5s 80us/step - loss: 0.0868 - acc: 0.9728
- val_loss: 0.0657 - val_acc: 0.9800
Epoch 16/20
60000/60000 [=====] - 5s 82us/step - loss: 0.0817 - acc: 0.9742
- val_loss: 0.0642 - val_acc: 0.9801
Epoch 17/20
60000/60000 [=====] - 5s 82us/step - loss: 0.0822 - acc: 0.9747
- val_loss: 0.0659 - val_acc: 0.9807
Epoch 18/20
60000/60000 [=====] - 5s 79us/step - loss: 0.0752 - acc: 0.9763
- val_loss: 0.0646 - val_acc: 0.9815
Epoch 19/20
60000/60000 [=====] - 5s 80us/step - loss: 0.0760 - acc: 0.9762
- val_loss: 0.0629 - val_acc: 0.9817
Epoch 20/20
60000/60000 [=====] - 5s 81us/step - loss: 0.0754 - acc: 0.9763
```

- val\_loss: 0.0604 -val\_acc: 0.9821

In [0]:

```
model1.summary()
```

Layer (type)	Output Shape	Param #
dense_15 (Dense)	(None, 256)	200960
batch_normalization_4 (Batch Normalization)	(None, 256)	1024
dropout_3 (Dropout)	(None, 256)	0
dense_16 (Dense)	(None, 128)	32896
batch_normalization_5 (Batch Normalization)	(None, 128)	512
dropout_4 (Dropout)	(None, 128)	0
dense_17 (Dense)	(None, 10)	1290
Total params: 236,682		
Trainable params: 235,914		
Non-trainable params: 768		

In [0]:

```
score = model1.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs


vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.06043866992703406

Test accuracy: 0.9821







**Model is better than others but not upto the model with dropout.**

In [0]:


```
w_after = model1.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure(figsize=(10,7))
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



## MODEL 2 - 3 HIDDEN LAYERS

### 2.1 MODEL WITH 2 POWERS AS NUMBER OF HIDDEN LAYERS (RELU ACTIVATION)

In [0]:

```
model2 = Sequential()
```

```
model2.add(Dense(512, input_dim=input_dim, activation='relu'))
model2.add(Dense(256, input_dim=input_dim, activation='relu'))
model2.add(Dense(128, input_dim=input_dim, activation='relu'))
model2.add(Dense(10, input_dim=input_dim, activation='softmax'))
```

In [0]:

```
model2.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model2.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1,
, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 4s 70us/step - loss: 0.2289 - acc: 0.9326  
- val\_loss: 0.1105 - val\_acc: 0.9641

Epoch 2/20

60000/60000 [=====] - 3s 56us/step - loss: 0.0846 - acc: 0.9738  
- val\_loss: 0.0812 - val\_acc: 0.9746

Epoch 3/20

60000/60000 [=====] - 3s 58us/step - loss: 0.0536 - acc: 0.9833  
- val\_loss: 0.0849 - val\_acc: 0.9735

Epoch 4/20

60000/60000 [=====] - 3s 56us/step - loss: 0.0396 - acc: 0.9866  
- val\_loss: 0.0728 - val\_acc: 0.9790

Epoch 5/20

60000/60000 [=====] - 3s 57us/step - loss: 0.0300 - acc: 0.9900  
- val\_loss: 0.0752 - val\_acc: 0.9786

Epoch 6/20

60000/60000 [=====] - 3s 58us/step - loss: 0.0241 - acc: 0.9921  
- val\_loss: 0.0730 - val\_acc: 0.9802

Epoch 7/20

60000/60000 [=====] - 3s 56us/step - loss: 0.0227 - acc: 0.9927  
- val\_loss: 0.0901 - val\_acc: 0.9739

Epoch 8/20

60000/60000 [=====] - 3s 58us/step - loss: 0.0181 - acc: 0.9943  
- val\_loss: 0.0865 - val\_acc: 0.9785

Epoch 9/20

60000/60000 [=====] - 3s 57us/step - loss: 0.0186 - acc: 0.9937  
- val\_loss: 0.0745 - val\_acc: 0.9810

Epoch 10/20

60000/60000 [=====] - 4s 59us/step - loss: 0.0157 - acc: 0.9947  
- val\_loss: 0.0907 - val\_acc: 0.9778

Epoch 11/20

60000/60000 [=====] - 3s 56us/step - loss: 0.0119 - acc: 0.9960  
- val\_loss: 0.0836 - val\_acc: 0.9809

Epoch 12/20

60000/60000 [=====] - 4s 60us/step - loss: 0.0139 - acc: 0.9956  
- val\_loss: 0.0754 - val\_acc: 0.9827

Epoch 13/20

60000/60000 [=====] - 3s 55us/step - loss: 0.0127 - acc: 0.9962  
- val\_loss: 0.0837 - val\_acc: 0.9800

Epoch 14/20

60000/60000 [=====] - 3s 58us/step - loss: 0.0112 - acc: 0.9964  
- val\_loss: 0.0921 - val\_acc: 0.9808

Epoch 15/20

60000/60000 [=====] - 3s 57us/step - loss: 0.0124 - acc: 0.9963  
- val\_loss: 0.0728 - val\_acc: 0.9841

Epoch 16/20

60000/60000 [=====] - 3s 55us/step - loss: 0.0086 - acc: 0.9974  
- val\_loss: 0.0999 - val\_acc: 0.9789

Epoch 17/20

60000/60000 [=====] - 3s 58us/step - loss: 0.0101 - acc: 0.9969  
- val\_loss: 0.0962 - val\_acc: 0.9799

Epoch 18/20

60000/60000 [=====] - 3s 56us/step - loss: 0.0079 - acc: 0.9975  
- val\_loss: 0.0980 - val\_acc: 0.9807

Epoch 19/20

60000/60000 [=====] - 3s 57us/step - loss: 0.0097 - acc: 0.9970  
- val\_loss: 0.1004 - val\_acc: 0.9804

Epoch 20/20

60000/60000 [=====] - 3s 58us/step - loss: 0.0092 - acc: 0.9975

- val\_loss: 0.1088 -val\_acc: 0.9788

In [0]:

```
model2.summary()
```

Layer (type)	Output Shape	Param #
dense_22 (Dense)	(None, 512)	401920
dense_23 (Dense)	(None, 256)	131328
dense_24 (Dense)	(None, 128)	32896
dense_25 (Dense)	(None, 10)	1290
Total params: 567,434		
Trainable params: 567,434		
Non-trainable params: 0		

In [0]:

```
score = model2.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.10877159107371262

Test accuracy: 0.9788



The model has overfit.

In [0]:

```
w_after = model2.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)

fig = plt.figure(figsize=(10,7))
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='r')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



## 2.2 RELU ACTIVATION WITH DROPOUT VALUE 0.5

In [0]:

```
from keras.layers import Dropout
model2 = Sequential()
model2.add(Dense(512, input_dim=input_dim, activation='relu'))
model2.add(Dropout(0.5))
model2.add(Dense(256, input_dim=input_dim, activation='relu'))
```

```
model2.add(Dropout(0.5))
model2.add(Dense(128, input_dim=input_dim, activation='relu'))
model2.add(Dropout(0.5))
model2.add(Dense(10, input_dim=input_dim, activation='softmax'))
```

In [0]:

```
model2.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model2.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1,
                    validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 4s 74us/step - loss: 0.5487 - acc: 0.8282  
- val\_loss: 0.1492 - val\_acc: 0.9546

Epoch 2/20

60000/60000 [=====] - 4s 61us/step - loss: 0.2275 - acc: 0.9369  
- val\_loss: 0.1108 - val\_acc: 0.9671

Epoch 3/20

60000/60000 [=====] - 4s 60us/step - loss: 0.1762 - acc: 0.9501  
- val\_loss: 0.0952 - val\_acc: 0.9717

Epoch 4/20

60000/60000 [=====] - 4s 61us/step - loss: 0.1512 - acc: 0.9580  
- val\_loss: 0.0834 - val\_acc: 0.9756

Epoch 5/20

60000/60000 [=====] - 4s 59us/step - loss: 0.1337 - acc: 0.9623  
- val\_loss: 0.0789 - val\_acc: 0.9776

Epoch 6/20

60000/60000 [=====] - 4s 59us/step - loss: 0.1204 - acc: 0.9666  
- val\_loss: 0.0791 - val\_acc: 0.9762

Epoch 7/20

60000/60000 [=====] - 4s 59us/step - loss: 0.1140 - acc: 0.9673  
- val\_loss: 0.0703 - val\_acc: 0.9781

Epoch 8/20

60000/60000 [=====] - 4s 60us/step - loss: 0.1030 - acc: 0.9716  
- val\_loss: 0.0697 - val\_acc: 0.9805

Epoch 9/20

60000/60000 [=====] - 4s 59us/step - loss: 0.0950 - acc: 0.9732  
- val\_loss: 0.0741 - val\_acc: 0.9786

Epoch 10/20

60000/60000 [=====] - 4s 59us/step - loss: 0.0903 - acc: 0.9742  
- val\_loss: 0.0715 - val\_acc: 0.9801

Epoch 11/20

60000/60000 [=====] - 4s 58us/step - loss: 0.0857 - acc: 0.9755  
- val\_loss: 0.0719 - val\_acc: 0.9792

Epoch 12/20

60000/60000 [=====] - 4s 60us/step - loss: 0.0837 - acc: 0.9756  
- val\_loss: 0.0707 - val\_acc: 0.9813

Epoch 13/20

60000/60000 [=====] - 4s 60us/step - loss: 0.0760 - acc: 0.9778  
- val\_loss: 0.0647 - val\_acc: 0.9815

Epoch 14/20

60000/60000 [=====] - 4s 58us/step - loss: 0.0733 - acc: 0.9781  
- val\_loss: 0.0718 - val\_acc: 0.9815

Epoch 15/20

60000/60000 [=====] - 4s 61us/step - loss: 0.0715 - acc: 0.9790  
- val\_loss: 0.0678 - val\_acc: 0.9825

Epoch 16/20

60000/60000 [=====] - 4s 61us/step - loss: 0.0681 - acc: 0.9802  
- val\_loss: 0.0687 - val\_acc: 0.9819

Epoch 17/20

60000/60000 [=====] - 4s 61us/step - loss: 0.0668 - acc: 0.9808  
- val\_loss: 0.0671 - val\_acc: 0.9822

Epoch 18/20

60000/60000 [=====] - 4s 61us/step - loss: 0.0650 - acc: 0.9813  
- val\_loss: 0.0649 - val\_acc: 0.9846

Epoch 19/20

60000/60000 [=====] - 4s 58us/step - loss: 0.0613 - acc: 0.9827  
- val\_loss: 0.0696 - val\_acc: 0.9828

Epoch 20/20

60000/60000 [=====] - 4s 61us/step - loss: 0.0592 - acc: 0.9823

- val\_loss: 0.0665 -val\_acc: 0.9828

In [0]:

```
model2.summary()
```

Layer (type)	Output Shape	Param #
dense_26 (Dense)	(None, 512)	401920
dropout_5 (Dropout)	(None, 512)	0
dense_27 (Dense)	(None, 256)	131328
dropout_6 (Dropout)	(None, 256)	0
dense_28 (Dense)	(None, 128)	32896
dropout_7 (Dropout)	(None, 128)	0
dense_29 (Dense)	(None, 10)	1290
Total params: 567,434		
Trainable params: 567,434		
Non-trainable params: 0		

In [0]:

```
score = model2.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.06652374192810584

Test accuracy: 0.9828





**A little overfitting is seen at the final few epochs**

In [0]:

```
w_after = model2.get_weights()

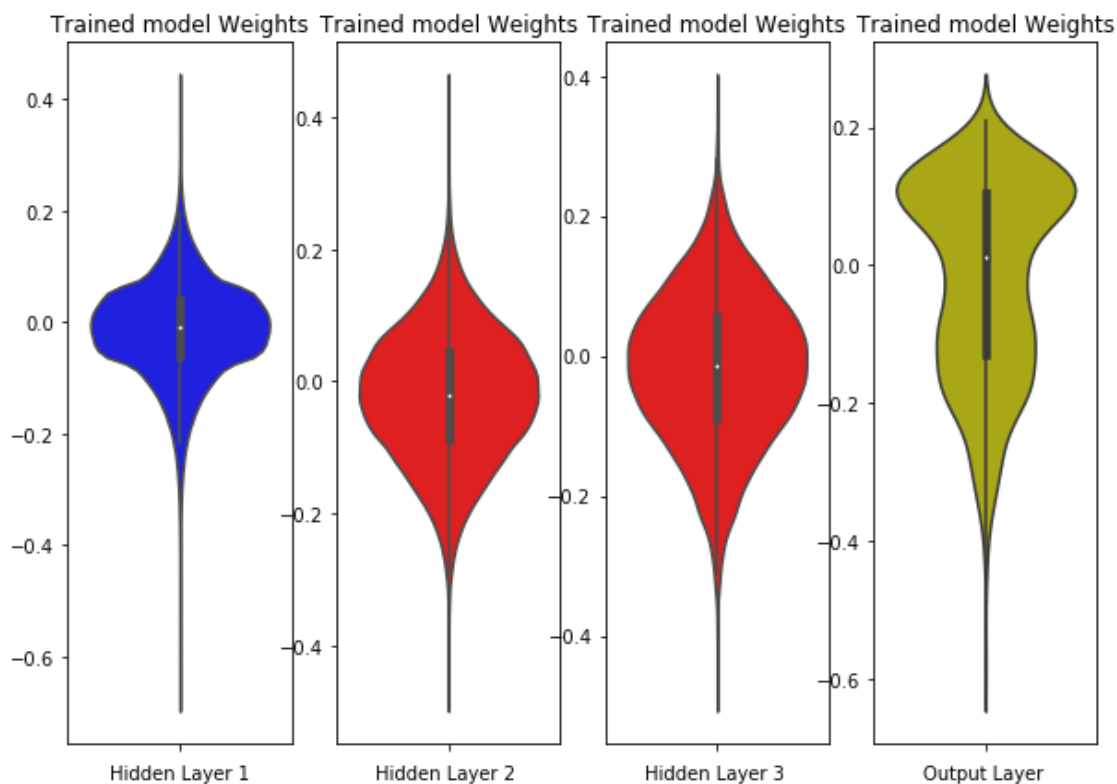
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)

fig = plt.figure(figsize=(10,7))
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='r')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



## 2.3 RELU ACTIVATION WITH BATCH NORM

In [0]:

```
from keras.layers.normalization import BatchNormalization
model2 = Sequential()
model2.add(Dense(512, input_dim=input_dim, activation='relu'))
model2.add(BatchNormalization())
model2.add(Dense(256, input_dim=input_dim, activation='relu'))
model2.add(BatchNormalization())
model2.add(Dense(128, input_dim=input_dim, activation='relu'))
model2.add(BatchNormalization())
model2.add(Dense(10, input_dim=input_dim, activation='softmax'))
```

In [0]:

```
model2.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model2.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1,
                    validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 8s 129us/step - loss: 0.1770 - acc: 0.9460
- val_loss: 0.1025 - val_acc: 0.9682
Epoch 2/20
60000/60000 [=====] - 6s 103us/step - loss: 0.0715 - acc: 0.9773
- val_loss: 0.0848 - val_acc: 0.9748
Epoch 3/20
60000/60000 [=====] - 6s 105us/step - loss: 0.0497 - acc: 0.9844
- val_loss: 0.0970 - val_acc: 0.9728
Epoch 4/20
60000/60000 [=====] - 6s 104us/step - loss: 0.0378 - acc: 0.9877
- val_loss: 0.0878 - val_acc: 0.9721
Epoch 5/20
60000/60000 [=====] - 6s 105us/step - loss: 0.0333 - acc: 0.9890
- val_loss: 0.0793 - val_acc: 0.9770
Epoch 6/20
60000/60000 [=====] - 6s 107us/step - loss: 0.0262 - acc: 0.9915
- val_loss: 0.0706 - val_acc: 0.9801
Epoch 7/20
60000/60000 [=====] - 6s 104us/step - loss: 0.0219 - acc: 0.9927
- val_loss: 0.0877 - val_acc: 0.9769
Epoch 8/20
60000/60000 [=====] - 6s 104us/step - loss: 0.0226 - acc: 0.9928
- val_loss: 0.0813 - val_acc: 0.9768
Epoch 9/20
60000/60000 [=====] - 6s 104us/step - loss: 0.0190 - acc: 0.9939
- val_loss: 0.0721 - val_acc: 0.9801
Epoch 10/20
60000/60000 [=====] - 6s 101us/step - loss: 0.0157 - acc: 0.9949
- val_loss: 0.0913 - val_acc: 0.9752
Epoch 11/20
60000/60000 [=====] - 6s 107us/step - loss: 0.0184 - acc: 0.9935
- val_loss: 0.0748 - val_acc: 0.9796
Epoch 12/20
60000/60000 [=====] - 6s 103us/step - loss: 0.0139 - acc: 0.9949
- val_loss: 0.0770 - val_acc: 0.9796
Epoch 13/20
60000/60000 [=====] - 6s 102us/step - loss: 0.0152 - acc: 0.9950
- val_loss: 0.0831 - val_acc: 0.9776
Epoch 14/20
60000/60000 [=====] - 6s 103us/step - loss: 0.0092 - acc: 0.9970
- val_loss: 0.0803 - val_acc: 0.9793
Epoch 15/20
60000/60000 [=====] - 6s 103us/step - loss: 0.0121 - acc: 0.9960
- val_loss: 0.0987 - val_acc: 0.9769
Epoch 16/20
60000/60000 [=====] - 6s 104us/step - loss: 0.0144 - acc: 0.9954
- val_loss: 0.0833 - val_acc: 0.9799
Epoch 17/20
60000/60000 [=====] - 6s 105us/step - loss: 0.0109 - acc: 0.9965
- val_loss: 0.0763 - val_acc: 0.9806
Epoch 18/20
60000/60000 [=====] - 6s 103us/step - loss: 0.0123 - acc: 0.9956
```



```
- val_loss: 0.0705 - val_acc: 0.9829
Epoch 19/20
60000/60000 [=====] - 6s 103us/step - loss: 0.0079 - acc: 0.9974
- val_loss: 0.0833 - val_acc: 0.9812
Epoch 20/20
60000/60000 [=====] - 6s 104us/step - loss: 0.0066 - acc: 0.9980
- val_loss: 0.0720 - val_acc: 0.9837
```

In [0]:

```
model2.summary()
```

Layer (type)	Output Shape	Param #
dense_30 (Dense)	(None, 512)	401920
batch_normalization_6 (Batch Normalization)	(None, 512)	2048
dense_31 (Dense)	(None, 256)	131328
batch_normalization_7 (Batch Normalization)	(None, 256)	1024
dense_32 (Dense)	(None, 128)	32896
batch_normalization_8 (Batch Normalization)	(None, 128)	512
dense_33 (Dense)	(None, 10)	1290
Total params: 571,018		
Trainable params: 569,226		
Non-trainable params: 1,792		

In [0]:

```
score = model2.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))


# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.07200342553501614  
Test accuracy: 0.9837





Again overfitting has occurred

In [0]:

```
w_after = model2.get_weights()

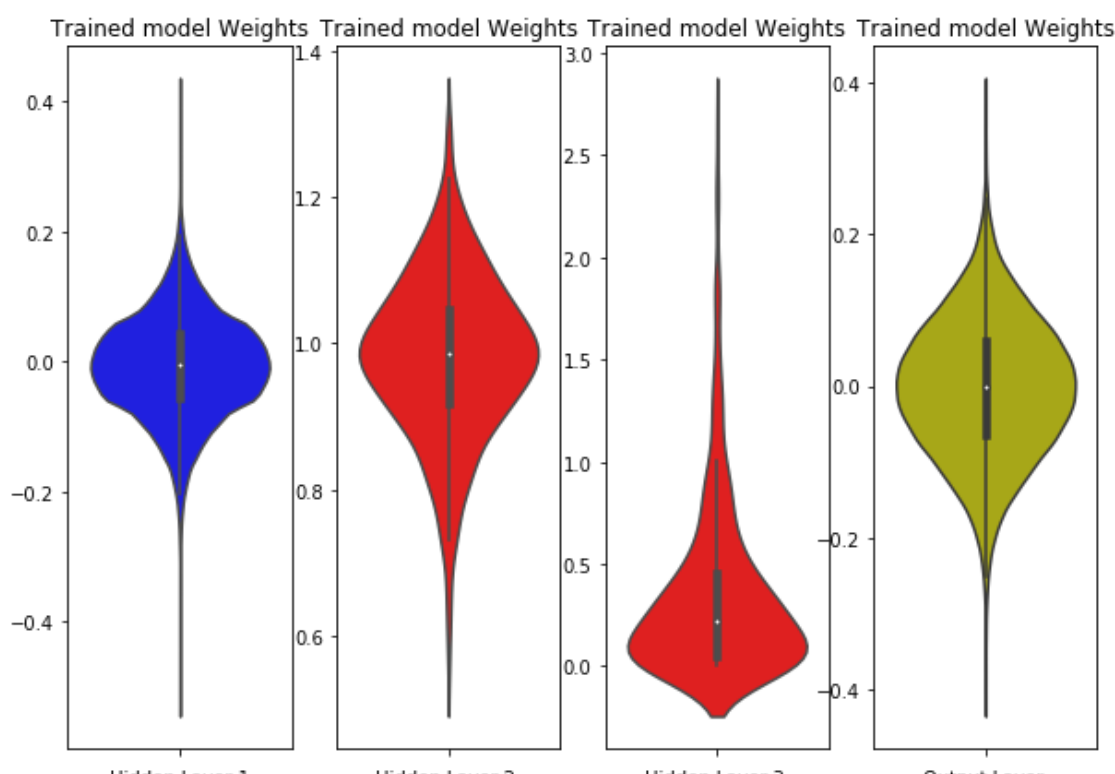
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)

fig = plt.figure(figsize=(10,7))
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='r')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



## 2.4 RELU ACTIVATION WITH BATCH NORM AND DROPOUT (0.5)

In [0]:

```
from keras.layers.normalization import BatchNormalization
model2 = Sequential()

model2.add(Dense(512, input_dim=input_dim, activation='relu'))
model2.add(BatchNormalization())
model2.add(Dropout(0.5))

model2.add(Dense(256, input_dim=input_dim, activation='relu'))
model2.add(BatchNormalization())
model2.add(Dropout(0.5))

model2.add(Dense(128, input_dim=input_dim, activation='relu'))
model2.add(BatchNormalization())
model2.add(Dropout(0.5))

model2.add(Dense(10, input_dim=input_dim, activation='softmax'))
```

In [0]:

```
model2.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model2.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1,
                    validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 8s 139us/step - loss: 0.5245 - acc: 0.8400  
- val\_loss: 0.1583 - val\_acc: 0.9524

Epoch 2/20

60000/60000 [=====] - 7s 110us/step - loss: 0.2361 - acc: 0.9308  
- val\_loss: 0.1084 - val\_acc: 0.9665

Epoch 3/20

60000/60000 [=====] - 6s 107us/step - loss: 0.1844 - acc: 0.9447  
- val\_loss: 0.0976 - val\_acc: 0.9702

Epoch 4/20

60000/60000 [=====] - 6s 108us/step - loss: 0.1586 - acc: 0.9538  
- val\_loss: 0.0869 - val\_acc: 0.9738

Epoch 5/20

60000/60000 [=====] - 7s 109us/step - loss: 0.1410 - acc: 0.9583  
- val\_loss: 0.0880 - val\_acc: 0.9741

Epoch 6/20

60000/60000 [=====] - 7s 109us/step - loss: 0.1284 - acc: 0.9615  
- val\_loss: 0.0775 - val\_acc: 0.9756

Epoch 7/20

60000/60000 [=====] - 6s 107us/step - loss: 0.1172 - acc: 0.9647  
- val\_loss: 0.0731 - val\_acc: 0.9770

Epoch 8/20

60000/60000 [=====] - 7s 111us/step - loss: 0.1089 - acc: 0.9678  
- val\_loss: 0.0630 - val\_acc: 0.9823

Epoch 9/20

60000/60000 [=====] - 7s 109us/step - loss: 0.1009 - acc: 0.9699  
- val\_loss: 0.0660 - val\_acc: 0.9812

Epoch 10/20

60000/60000 [=====] - 7s 109us/step - loss: 0.0976 - acc: 0.9708  
- val\_loss: 0.0692 - val\_acc: 0.9803

Epoch 11/20

60000/60000 [=====] - 6s 107us/step - loss: 0.0904 - acc: 0.9725  
- val\_loss: 0.0674 - val\_acc: 0.9810

Epoch 12/20

60000/60000 [=====] - 6s 108us/step - loss: 0.0887 - acc: 0.9731  
- val\_loss: 0.0657 - val\_acc: 0.9804

Epoch 13/20

60000/60000 [=====] - 7s 110us/step - loss: 0.0816 - acc: 0.9749  
- val\_loss: 0.0664 - val\_acc: 0.9812

Epoch 14/20

60000/60000 [=====] - 6s 106us/step - loss: 0.0801 - acc: 0.9754  
- val\_loss: 0.0664 - val\_acc: 0.9812

```

60000/60000 [=====] - 6s 106us/step - loss: 0.0801 - acc: 0.9754
- val_loss: 0.0564 - val_acc: 0.9831
Epoch 15/20
60000/60000 [=====] - 7s 110us/step - loss: 0.0749 - acc: 0.9774
- val_loss: 0.0570 - val_acc: 0.9837
Epoch 16/20
60000/60000 [=====] - 7s 112us/step - loss: 0.0706 - acc: 0.9784
- val_loss: 0.0624 - val_acc: 0.9816
Epoch 17/20
60000/60000 [=====] - 6s 108us/step - loss: 0.0727 - acc: 0.9782
- val_loss: 0.0621 - val_acc: 0.9835
Epoch 18/20
60000/60000 [=====] - 6s 107us/step - loss: 0.0691 - acc: 0.9781
- val_loss: 0.0628 - val_acc: 0.9822
Epoch 19/20
60000/60000 [=====] - 7s 109us/step - loss: 0.0636 - acc: 0.9801
- val_loss: 0.0622 - val_acc: 0.9824
Epoch 20/20
60000/60000 [=====] - 7s 110us/step - loss: 0.0632 - acc: 0.9803
- val_loss: 0.0571 - val_acc: 0.9844

```

In [0]:

```
model2.summary()
```

Layer (type)	Output Shape	Param #
dense_38 (Dense)	(None, 512)	401920
batch_normalization_12 (Batch Normalization)	(None, 512)	2048
dropout_11 (Dropout)	(None, 512)	0
dense_39 (Dense)	(None, 256)	131328
batch_normalization_13 (Batch Normalization)	(None, 256)	1024
dropout_12 (Dropout)	(None, 256)	0
dense_40 (Dense)	(None, 128)	32896
batch_normalization_14 (Batch Normalization)	(None, 128)	512
dropout_13 (Dropout)	(None, 128)	0
dense_41 (Dense)	(None, 10)	1290
Total params: 571,018		
Trainable params: 569,226		
Non-trainable params: 1,792		

In [0]:

```

score = model2.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss

```


```
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.057054539459303485

Test accuracy: 0.9844



**The best model in 3 hidden layer models.**

In [0]:

```
w_after = model2.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)

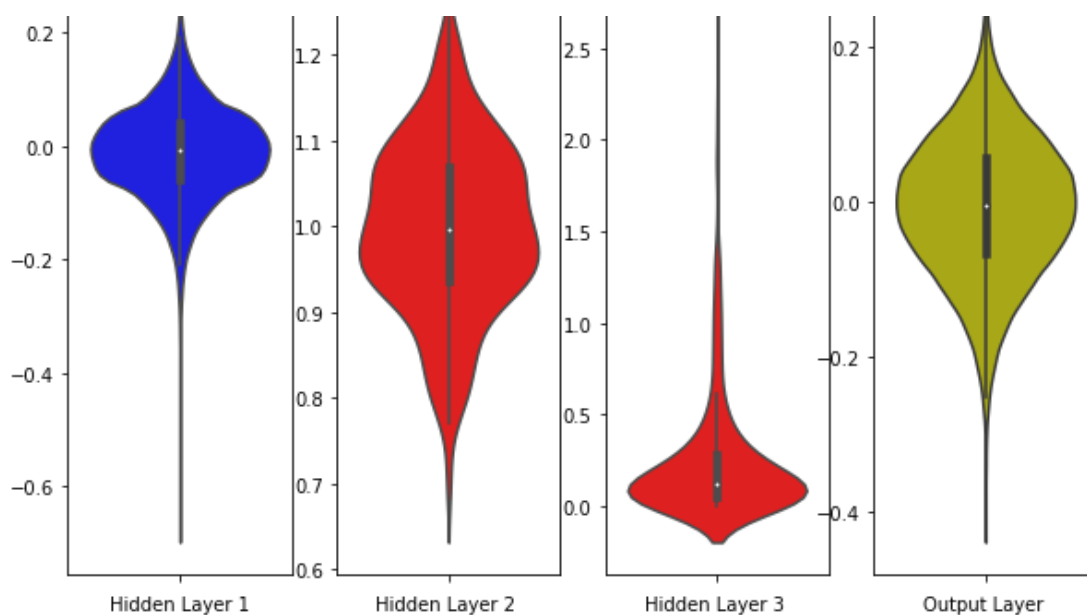
fig = plt.figure(figsize=(10,7))
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='r')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```





## MODEL 3 - 5 HIDDEN LAYERS

### 3.1 MODEL WITH 2 POWERS AS NUMBER OF HIDDEN LAYERS (RELU ACTIVATION)

In [0]:

```
model3 = Sequential()
model3.add(Dense(512, input_dim=input_dim, activation='relu'))
model3.add(Dense(256, input_dim=input_dim, activation='relu'))
model3.add(Dense(128, input_dim=input_dim, activation='relu'))
model3.add(Dense(64, input_dim=input_dim, activation='relu'))
model3.add(Dense(10, input_dim=input_dim, activation='softmax'))
```

In [0]:

```
model3.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model3.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1,
                    validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 5s 85us/step - loss: 0.2399 - acc: 0.9279  
- val\_loss: 0.1218 - val\_acc: 0.9621

Epoch 2/20

60000/60000 [=====] - 4s 63us/step - loss: 0.0875 - acc: 0.9728  
- val\_loss: 0.0844 - val\_acc: 0.9743

Epoch 3/20

60000/60000 [=====] - 4s 62us/step - loss: 0.0589 - acc: 0.9813  
- val\_loss: 0.0850 - val\_acc: 0.9735

Epoch 4/20

60000/60000 [=====] - 4s 63us/step - loss: 0.0432 - acc: 0.9866  
- val\_loss: 0.0687 - val\_acc: 0.9798

Epoch 5/20

60000/60000 [=====] - 4s 62us/step - loss: 0.0329 - acc: 0.9889  
- val\_loss: 0.0798 - val\_acc: 0.9769

Epoch 6/20

60000/60000 [=====] - 4s 61us/step - loss: 0.0307 - acc: 0.9899  
- val\_loss: 0.0676 - val\_acc: 0.9804

Epoch 7/20

60000/60000 [=====] - 4s 61us/step - loss: 0.0249 - acc: 0.9920  
- val\_loss: 0.0768 - val\_acc: 0.9802

Epoch 8/20

60000/60000 [=====] - 4s 61us/step - loss: 0.0216 - acc: 0.9931  
- val\_loss: 0.0796 - val\_acc: 0.9792

Epoch 9/20

60000/60000 [=====] - 4s 62us/step - loss: 0.0197 - acc: 0.9935  
- val\_loss: 0.0846 - val\_acc: 0.9796

Epoch 10/20

```
Epoch 10/20
60000/60000 [=====] - 4s 63us/step - loss: 0.0177 - acc: 0.9940
- val_loss: 0.0771 - val_acc: 0.9803
Epoch 11/20
60000/60000 [=====] - 4s 62us/step - loss: 0.0142 - acc: 0.9955
- val_loss: 0.0869 - val_acc: 0.9793
Epoch 12/20
60000/60000 [=====] - 4s 63us/step - loss: 0.0160 - acc: 0.9951
- val_loss: 0.0756 - val_acc: 0.9815
Epoch 13/20
60000/60000 [=====] - 4s 63us/step - loss: 0.0145 - acc: 0.9955
- val_loss: 0.0874 - val_acc: 0.9786
Epoch 14/20
60000/60000 [=====] - 4s 62us/step - loss: 0.0124 - acc: 0.9960
- val_loss: 0.0932 - val_acc: 0.9793
Epoch 15/20
60000/60000 [=====] - 4s 62us/step - loss: 0.0129 - acc: 0.9960
- val_loss: 0.0936 - val_acc: 0.9807
Epoch 16/20
60000/60000 [=====] - 4s 62us/step - loss: 0.0115 - acc: 0.9963
- val_loss: 0.0906 - val_acc: 0.9785
Epoch 17/20
60000/60000 [=====] - 4s 62us/step - loss: 0.0088 - acc: 0.9974
- val_loss: 0.1056 - val_acc: 0.9775
Epoch 18/20
60000/60000 [=====] - 4s 63us/step - loss: 0.0135 - acc: 0.9957
- val_loss: 0.0892 - val_acc: 0.9811
Epoch 19/20
60000/60000 [=====] - 4s 62us/step - loss: 0.0072 - acc: 0.9979
- val_loss: 0.1064 - val_acc: 0.9784
Epoch 20/20
60000/60000 [=====] - 4s 62us/step - loss: 0.0112 - acc: 0.9965
- val_loss: 0.0849 - val_acc: 0.9813
```

In [0]:

```
model3.summary()
```

Layer (type)	Output Shape	Param #
dense_42 (Dense)	(None, 512)	401920
dense_43 (Dense)	(None, 256)	131328
dense_44 (Dense)	(None, 128)	32896
dense_45 (Dense)	(None, 64)	8256
dense_46 (Dense)	(None, 10)	650
Total params: 575,050		
Trainable params: 575,050		
Non-trainable params: 0		

In [0]:

```
score = model3.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```


```
# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.08486868686463067

Test accuracy: 0.9813



**Overfitting is observed.**

In [0]:

```
w_after = model3.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
out_w = w_after[8].flatten().reshape(-1,1)

fig = plt.figure(figsize=(10,7))
plt.title("Weight matrices after model trained")
plt.subplot(1, 5, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 5, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

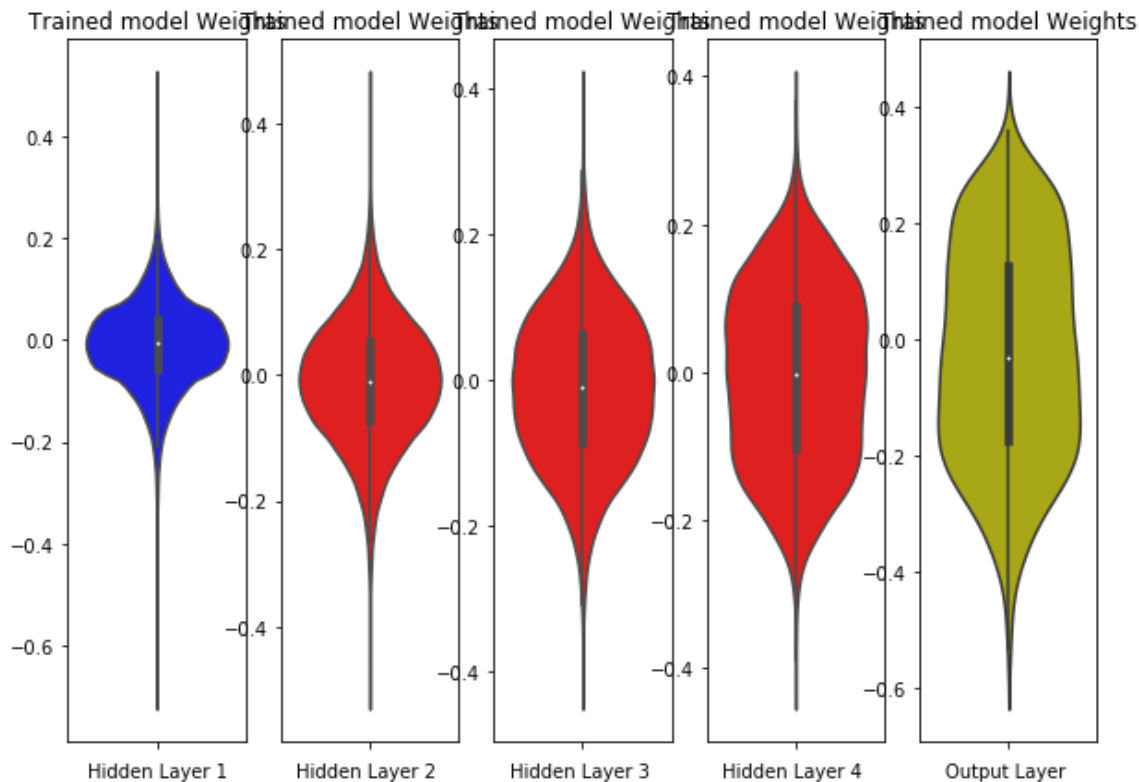
plt.subplot(1, 5, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='r')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 5, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w, color='r')
plt.xlabel('Hidden Layer 4 ')

plt.subplot(1, 5, 5)
plt.title("Trained model Weights")
```



```
ax = sns.violinplot(y=out_w, color='y')
plt.xlabel('Output Layer ')
plt.show()
```



### 3.2 RELU ACTIVATION WITH DROPOUT VALUE 0.5

In [0]:

```
from keras.layers import Dropout
model3 = Sequential()
model3.add(Dense(512, input_dim=input_dim, activation='relu'))
model3.add(Dropout(0.5))
model3.add(Dense(256, input_dim=input_dim, activation='relu'))
model3.add(Dropout(0.5))
model3.add(Dense(128, input_dim=input_dim, activation='relu'))
model3.add(Dropout(0.5))
model3.add(Dense(64, input_dim=input_dim, activation='relu'))
model3.add(Dropout(0.5))
model3.add(Dense(10, input_dim=input_dim, activation='softmax'))
```

In [0]:

```
model3.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model3.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1,
                    validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 6s 95us/step - loss: 0.8366 - acc: 0.7285
- val_loss: 0.2140 - val_acc: 0.9408
Epoch 2/20
60000/60000 [=====] - 4s 65us/step - loss: 0.3213 - acc: 0.9203
- val_loss: 0.1466 - val_acc: 0.9608
Epoch 3/20
60000/60000 [=====] - 4s 66us/step - loss: 0.2444 - acc: 0.9390
- val_loss: 0.1272 - val_acc: 0.9682
Epoch 4/20
60000/60000 [=====] - 4s 65us/step - loss: 0.2060 - acc: 0.9502
- val_loss: 0.1079 - val_acc: 0.9716
Epoch 5/20
60000/60000 [=====] - 4s 66us/step - loss: 0.1863 - acc: 0.9544
- val_loss: 0.1049 - val_acc: 0.9723
Epoch 6/20
60000/60000 [=====] - 4s 67us/step - loss: 0.1709 - acc: 0.9578
```

```

- val_loss: 0.1066 - val_acc: 0.9739
Epoch 7/20
60000/60000 [=====] - 4s 66us/step - loss: 0.1585 - acc: 0.9613
- val_loss: 0.1017 - val_acc: 0.9738
Epoch 8/20
60000/60000 [=====] - 4s 67us/step - loss: 0.1455 - acc: 0.9637
- val_loss: 0.0901 - val_acc: 0.9760
Epoch 9/20
60000/60000 [=====] - 4s 66us/step - loss: 0.1379 - acc: 0.9660
- val_loss: 0.0914 - val_acc: 0.9772
Epoch 10/20
60000/60000 [=====] - 4s 68us/step - loss: 0.1343 - acc: 0.9676
- val_loss: 0.0839 - val_acc: 0.9777
Epoch 11/20
60000/60000 [=====] - 4s 66us/step - loss: 0.1212 - acc: 0.9697
- val_loss: 0.0862 - val_acc: 0.9787
Epoch 12/20
60000/60000 [=====] - 4s 66us/step - loss: 0.1175 - acc: 0.9718
- val_loss: 0.0805 - val_acc: 0.9799
Epoch 13/20
60000/60000 [=====] - 4s 66us/step - loss: 0.1113 - acc: 0.9734
- val_loss: 0.0878 - val_acc: 0.9772
Epoch 14/20
60000/60000 [=====] - 4s 65us/step - loss: 0.1093 - acc: 0.9728
- val_loss: 0.0852 - val_acc: 0.9789
Epoch 15/20
60000/60000 [=====] - 4s 67us/step - loss: 0.1058 - acc: 0.9742
- val_loss: 0.0893 - val_acc: 0.9791
Epoch 16/20
60000/60000 [=====] - 4s 67us/step - loss: 0.1044 - acc: 0.9741
- val_loss: 0.0881 - val_acc: 0.9792
Epoch 17/20
60000/60000 [=====] - 4s 66us/step - loss: 0.0997 - acc: 0.9752
- val_loss: 0.0779 - val_acc: 0.9807
Epoch 18/20
60000/60000 [=====] - 4s 67us/step - loss: 0.0953 - acc: 0.9756
- val_loss: 0.0801 - val_acc: 0.9814
Epoch 19/20
60000/60000 [=====] - 4s 66us/step - loss: 0.0911 - acc: 0.9772
- val_loss: 0.0870 - val_acc: 0.9812
Epoch 20/20
60000/60000 [=====] - 4s 66us/step - loss: 0.0961 - acc: 0.9763
- val_loss: 0.0736 - val_acc: 0.9820

```

In [0]:

```
model3.summary()
```

Layer (type)	Output Shape	Param #
=====		
dense_47 (Dense)	(None, 512)	401920
dropout_14 (Dropout)	(None, 512)	0
dense_48 (Dense)	(None, 256)	131328
dropout_15 (Dropout)	(None, 256)	0
dense_49 (Dense)	(None, 128)	32896
dropout_16 (Dropout)	(None, 128)	0
dense_50 (Dense)	(None, 64)	8256
dropout_17 (Dropout)	(None, 64)	0
dense_51 (Dense)	(None, 10)	650
=====		
Total params: 575,050		
Trainable params: 575,050		
Non-trainable params: 0		

In [0]:

```
score = model3.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.07358933680157234

Test accuracy: 0.982



**The best model in 5 hidden layer models.**

In [0]:

```
w_after = model3.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
out_w = w_after[8].flatten().reshape(-1,1)

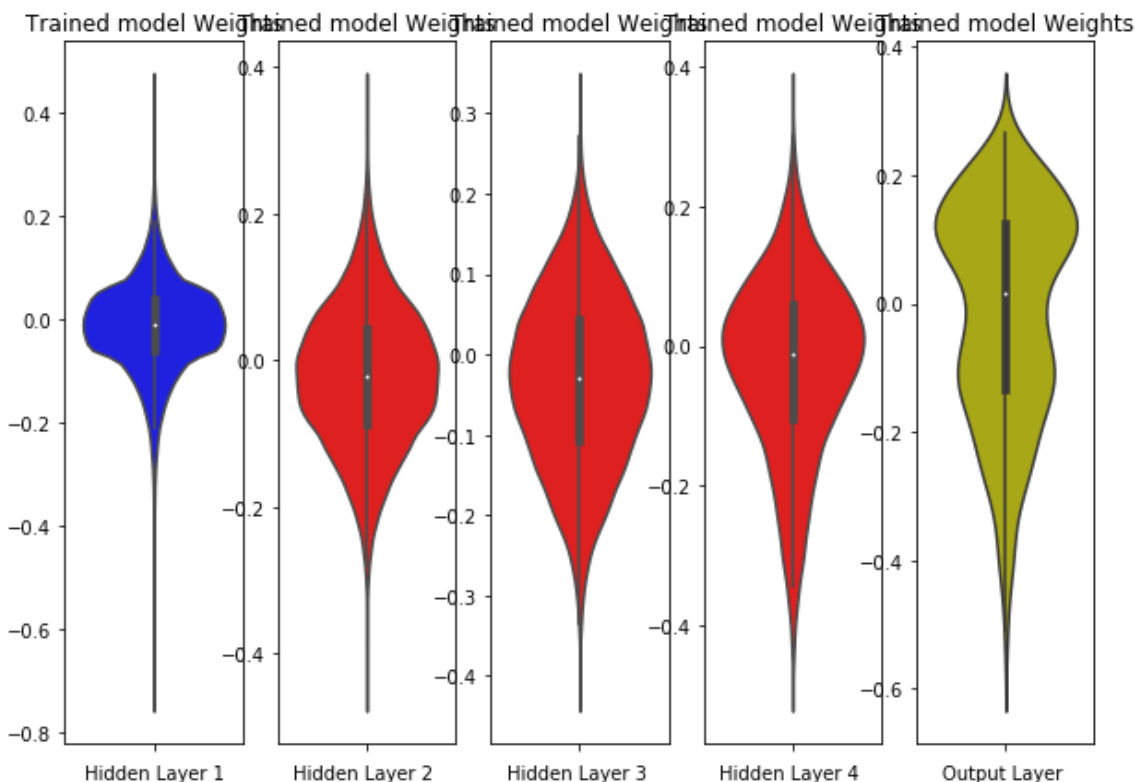
fig = plt.figure(figsize=(10,7))
plt.title("Weight matrices after model trained")
plt.subplot(1, 5, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w, color='b')
plt.xlabel('Hidden Layer 1')
```

```
plt.subplot(1, 5, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 5, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='r')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 5, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w, color='r')
plt.xlabel('Hidden Layer 4 ')

plt.subplot(1, 5, 5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w, color='y')
plt.xlabel('Output Layer ')
plt.show()
```



### 3.3 RELU ACTIVATION WITH BATCH NORM

In [0]:

```
from keras.layers.normalization import BatchNormalization
model3 = Sequential()
model3.add(Dense(512, input_dim=input_dim, activation='relu'))
model3.add(BatchNormalization())
model3.add(Dense(256, input_dim=input_dim, activation='relu'))
model3.add(BatchNormalization())
model3.add(Dense(128, input_dim=input_dim, activation='relu'))
model3.add(BatchNormalization())
model3.add(Dense(64, input_dim=input_dim, activation='relu'))
model3.add(BatchNormalization())
model3.add(Dense(10, input_dim=input_dim, activation='softmax'))
```

In [0]:

```
model3.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model3.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1,
                    validation_data=(X_test, Y_test))
```

```

Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 10s 162us/step - loss: 0.2005 - acc: 0.9407 - val_loss: 0.0970 - val_acc: 0.9689
Epoch 2/20
60000/60000 [=====] - 7s 124us/step - loss: 0.0809 - acc: 0.9753 - val_loss: 0.1065 - val_acc: 0.9652
Epoch 3/20
60000/60000 [=====] - 8s 126us/step - loss: 0.0576 - acc: 0.9820 - val_loss: 0.1053 - val_acc: 0.9676
Epoch 4/20
60000/60000 [=====] - 8s 127us/step - loss: 0.0441 - acc: 0.9858 - val_loss: 0.0779 - val_acc: 0.9763
Epoch 5/20
60000/60000 [=====] - 7s 122us/step - loss: 0.0354 - acc: 0.9886 - val_loss: 0.0798 - val_acc: 0.9771
Epoch 6/20
60000/60000 [=====] - 8s 126us/step - loss: 0.0313 - acc: 0.9899 - val_loss: 0.0671 - val_acc: 0.9819
Epoch 7/20
60000/60000 [=====] - 7s 122us/step - loss: 0.0272 - acc: 0.9911 - val_loss: 0.0701 - val_acc: 0.9792
Epoch 8/20
60000/60000 [=====] - 7s 123us/step - loss: 0.0254 - acc: 0.9913 - val_loss: 0.0842 - val_acc: 0.9780
Epoch 9/20
60000/60000 [=====] - 7s 123us/step - loss: 0.0220 - acc: 0.9925 - val_loss: 0.0711 - val_acc: 0.9792
Epoch 10/20
60000/60000 [=====] - 7s 123us/step - loss: 0.0212 - acc: 0.9930 - val_loss: 0.0845 - val_acc: 0.9770
Epoch 11/20
60000/60000 [=====] - 7s 122us/step - loss: 0.0192 - acc: 0.9936 - val_loss: 0.0822 - val_acc: 0.9788
Epoch 12/20
60000/60000 [=====] - 8s 126us/step - loss: 0.0150 - acc: 0.9952 - val_loss: 0.0762 - val_acc: 0.9803
Epoch 13/20
60000/60000 [=====] - 7s 121us/step - loss: 0.0172 - acc: 0.9943 - val_loss: 0.0726 - val_acc: 0.9804
Epoch 14/20
60000/60000 [=====] - 7s 124us/step - loss: 0.0181 - acc: 0.9940 - val_loss: 0.0798 - val_acc: 0.9786
Epoch 15/20
60000/60000 [=====] - 7s 123us/step - loss: 0.0126 - acc: 0.9957 - val_loss: 0.0720 - val_acc: 0.9809
Epoch 16/20
60000/60000 [=====] - 8s 127us/step - loss: 0.0119 - acc: 0.9959 - val_loss: 0.0807 - val_acc: 0.9803
Epoch 17/20
60000/60000 [=====] - 7s 124us/step - loss: 0.0120 - acc: 0.9959 - val_loss: 0.0796 - val_acc: 0.9814
Epoch 18/20
60000/60000 [=====] - 7s 123us/step - loss: 0.0134 - acc: 0.9955 - val_loss: 0.0792 - val_acc: 0.9802
Epoch 19/20
60000/60000 [=====] - 7s 124us/step - loss: 0.0104 - acc: 0.9965 - val_loss: 0.0752 - val_acc: 0.9820
Epoch 20/20
60000/60000 [=====] - 7s 123us/step - loss: 0.0108 - acc: 0.9963 - val_loss: 0.1005 - val_acc: 0.9774

```

In [0]:

```
model3.summary()
```

Layer (type)	Output Shape	Param #
dense_52 (Dense)	(None, 512)	401920
batch normalization_15 (Batch Normalization)	(None, 512)	2048

dense_53 (Dense)	(None, 256)	131328
batch_normalization_16 (Batch Normalization)	(None, 256)	1024
dense_54 (Dense)	(None, 128)	32896
batch_normalization_17 (Batch Normalization)	(None, 128)	512
dense_55 (Dense)	(None, 64)	8256
batch_normalization_18 (Batch Normalization)	(None, 64)	256
dense_56 (Dense)	(None, 10)	650
=====		
Total params: 578,890		
Trainable params: 576,970		
Non-trainable params: 1,920		

In [0]:

```
score = model3.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.1004731001386419

Test accuracy: 0.9774



Overfitting at the very early epochs.

```
In [0]:
```

```
w_after = model3.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
out_w = w_after[8].flatten().reshape(-1,1)

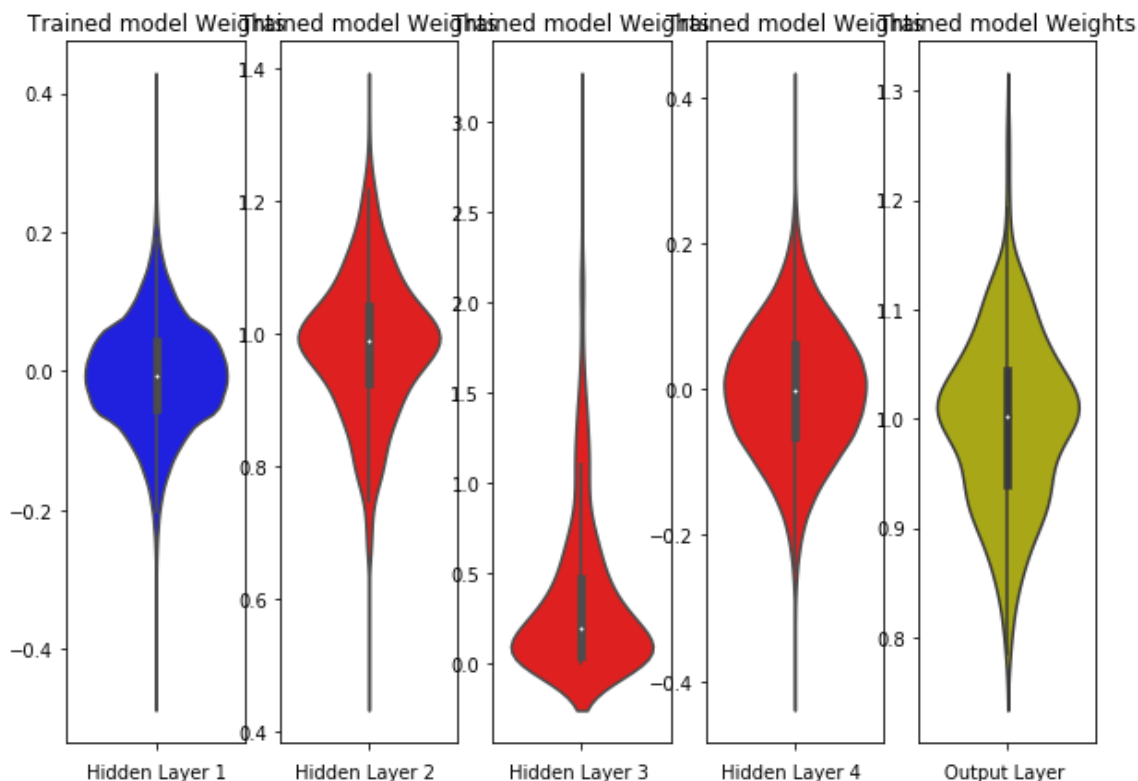
fig = plt.figure(figsize=(10,7))
plt.title("Weight matrices after model trained")
plt.subplot(1, 5, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 5, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 5, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='r')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 5, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w, color='r')
plt.xlabel('Hidden Layer 4 ')

plt.subplot(1, 5, 5)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



### 3.4 RELU ACTIVATION WITH BATCH NORM AND DROPOUT (0.5)

```
In [0]:
```

```
from keras.layers.normalization import BatchNormalization
```

```

model3 = Sequential()

model3.add(Dense(512, input_dim=input_dim, activation='relu'))
model3.add(BatchNormalization())
model3.add(Dropout(0.5))

model3.add(Dense(256, input_dim=input_dim, activation='relu'))
model3.add(BatchNormalization())
model3.add(Dropout(0.5))

model3.add(Dense(128, input_dim=input_dim, activation='relu'))
model3.add(BatchNormalization())
model3.add(Dropout(0.5))

model3.add(Dense(64, input_dim=input_dim, activation='relu'))
model3.add(BatchNormalization())
model3.add(Dropout(0.5))

model3.add(Dense(10, input_dim=input_dim, activation='softmax'))

```

In [0]:

```

model3.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model3.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1,
    validation_data=(X_test, Y_test))

```

Train on 60000 samples, validate on 10000 samples

```

Epoch 1/20
60000/60000 [=====] - 11s 180us/step - loss: 0.8112 - acc: 0.753
4 - val_loss: 0.2031 - val_acc: 0.9414
Epoch 2/20
60000/60000 [=====] - 8s 131us/step - loss: 0.3121 - acc: 0.9153
- val_loss: 0.1321 - val_acc: 0.9625
Epoch 3/20
60000/60000 [=====] - 8s 129us/step - loss: 0.2375 - acc: 0.9369
- val_loss: 0.1194 - val_acc: 0.9650
Epoch 4/20
60000/60000 [=====] - 8s 129us/step - loss: 0.2027 - acc: 0.9460
- val_loss: 0.1066 - val_acc: 0.9689
Epoch 5/20
60000/60000 [=====] - 8s 130us/step - loss: 0.1734 - acc: 0.9541
- val_loss: 0.0918 - val_acc: 0.9746
Epoch 6/20
60000/60000 [=====] - 8s 131us/step - loss: 0.1629 - acc: 0.9556
- val_loss: 0.0926 - val_acc: 0.9748
Epoch 7/20
60000/60000 [=====] - 8s 128us/step - loss: 0.1523 - acc: 0.9597
- val_loss: 0.0850 - val_acc: 0.9765
Epoch 8/20
60000/60000 [=====] - 8s 130us/step - loss: 0.1430 - acc: 0.9621
- val_loss: 0.0874 - val_acc: 0.9748
Epoch 9/20
60000/60000 [=====] - 8s 130us/step - loss: 0.1335 - acc: 0.9647
- val_loss: 0.0760 - val_acc: 0.9789
Epoch 10/20
60000/60000 [=====] - 8s 131us/step - loss: 0.1239 - acc: 0.9666
- val_loss: 0.0757 - val_acc: 0.9774
Epoch 11/20
60000/60000 [=====] - 8s 130us/step - loss: 0.1204 - acc: 0.9673
- val_loss: 0.0717 - val_acc: 0.9809
Epoch 12/20
60000/60000 [=====] - 8s 128us/step - loss: 0.1138 - acc: 0.9685
- val_loss: 0.0737 - val_acc: 0.9812
Epoch 13/20
60000/60000 [=====] - 8s 133us/step - loss: 0.1081 - acc: 0.9719
- val_loss: 0.0699 - val_acc: 0.9823
Epoch 14/20
60000/60000 [=====] - 8s 131us/step - loss: 0.1038 - acc: 0.9720
- val_loss: 0.0680 - val_acc: 0.9828
Epoch 15/20
60000/60000 [=====] - 8s 129us/step - loss: 0.1028 - acc: 0.9718

```



```

- val_loss: 0.0694 - val_acc: 0.9810
Epoch 16/20
60000/60000 [=====] - 8s 134us/step - loss: 0.0914 - acc: 0.9752
- val_loss: 0.0673 - val_acc: 0.9819
Epoch 17/20
60000/60000 [=====] - 8s 132us/step - loss: 0.0956 - acc: 0.9747
- val_loss: 0.0631 - val_acc: 0.9829
Epoch 18/20
60000/60000 [=====] - 8s 130us/step - loss: 0.0907 - acc: 0.9757
- val_loss: 0.0654 - val_acc: 0.9825
Epoch 19/20
60000/60000 [=====] - 8s 130us/step - loss: 0.0867 - acc: 0.9762
- val_loss: 0.0676 - val_acc: 0.9827
Epoch 20/20
60000/60000 [=====] - 8s 130us/step - loss: 0.0820 - acc: 0.9777
- val_loss: 0.0676 - val_acc: 0.9821

```

In [0]:

```
model3.summary()
```

Layer (type)	Output Shape	Param #
dense_57 (Dense)	(None, 512)	401920
batch_normalization_19 (Batch Normalization)	(None, 512)	2048
dropout_18 (Dropout)	(None, 512)	0
dense_58 (Dense)	(None, 256)	131328
batch_normalization_20 (Batch Normalization)	(None, 256)	1024
dropout_19 (Dropout)	(None, 256)	0
dense_59 (Dense)	(None, 128)	32896
batch_normalization_21 (Batch Normalization)	(None, 128)	512
dropout_20 (Dropout)	(None, 128)	0
dense_60 (Dense)	(None, 64)	8256
batch_normalization_22 (Batch Normalization)	(None, 64)	256
dropout_21 (Dropout)	(None, 64)	0
dense_61 (Dense)	(None, 10)	650
Total params: 578,890		
Trainable params: 576,970		
Non-trainable params: 1,920		

In [0]:

```

score = model3.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

```


```
# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.06762448588523548

Test accuracy: 0.9821



**Better than many models.**

In [0]:

```
w_after = model3.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
h4_w = w_after[6].flatten().reshape(-1,1)
out_w = w_after[8].flatten().reshape(-1,1)

fig = plt.figure(figsize=(10,7))
plt.title("Weight matrices after model trained")
plt.subplot(1, 5, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

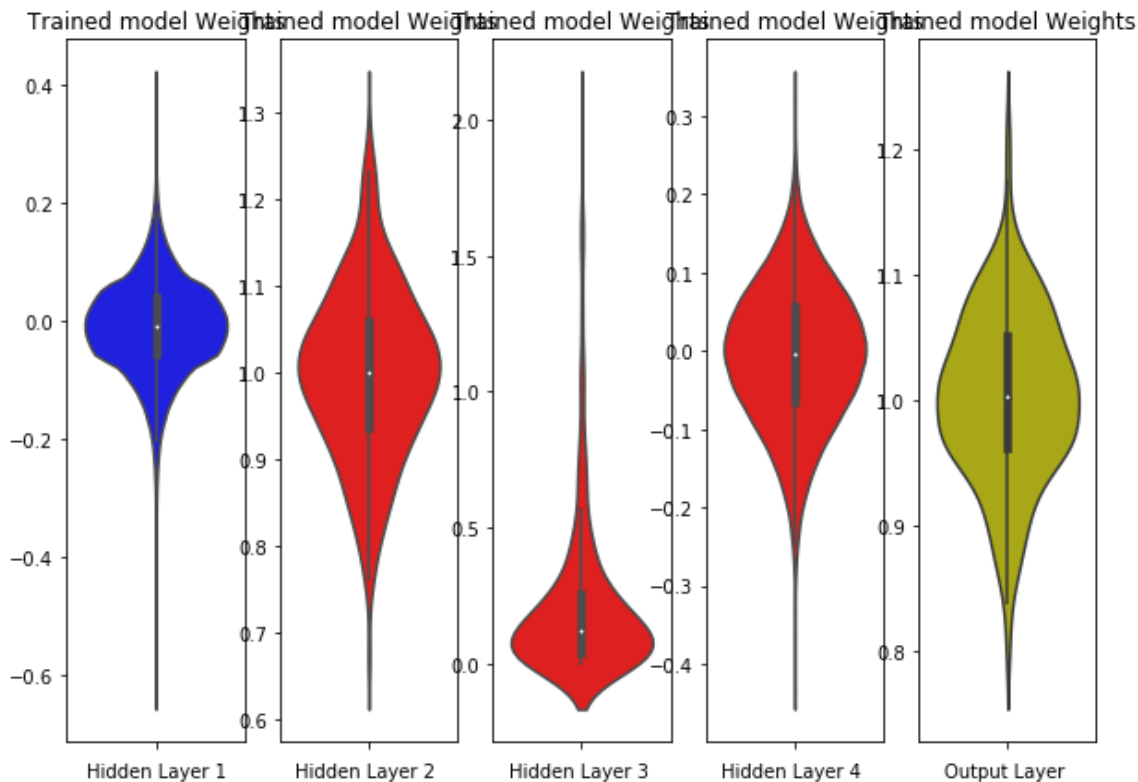
plt.subplot(1, 5, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 5, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='r')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 5, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h4_w, color='r')
plt.xlabel('Hidden Layer 4 ')

plt.subplot(1, 5, 5)
plt.title("Trained model Weights")
```

```
ax = sns.violinplot(y=out_w, color='y')
plt.xlabel('Output Layer')
plt.show()
```



#### 4.1 RELU ACTIVATION WITH DROPOUT VALUE 0.5 (Optimiser - RMSPROP)

In [0]:

```
from keras.layers import Dropout
model4 = Sequential()
model4.add(Dense(512, input_dim=input_dim, activation='relu'))
model4.add(Dropout(0.5))
model4.add(Dense(256, input_dim=input_dim, activation='relu'))
model4.add(Dropout(0.5))
model4.add(Dense(128, input_dim=input_dim, activation='relu'))
model4.add(Dropout(0.5))
model4.add(Dense(10, input_dim=input_dim, activation='softmax'))
```

In [15]:

```
model4.compile(optimizer='RMSprop', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model4.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1,
                    validation_data=(X_test, Y_test))
```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/python/ops/math\_grad.py:1250: add\_dispatch\_support.<locals>.wrapper (from tensorflow.python.ops.array\_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Use tf.where in 2.0, which has the same broadcast rule as np.where

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 8s 134us/step - loss: 0.5014 - acc: 0.8458

- val\_loss: 0.1567 - val\_acc: 0.9526

Epoch 2/20

60000/60000 [=====] - 3s 54us/step - loss: 0.2250 - acc: 0.9398

- val\_loss: 0.1244 - val\_acc: 0.9663

Epoch 3/20

60000/60000 [=====] - 3s 53us/step - loss: 0.1815 - acc: 0.9515

- val\_loss: 0.1005 - val\_acc: 0.9705

Epoch 4/20

60000/60000 [=====] - 3s 54us/step - loss: 0.1602 - acc: 0.9590

- val\_loss: 0.1014 - val\_acc: 0.9744

Epoch 5/20

```

Epoch 6/20
60000/60000 [=====] - 3s 53us/step - loss: 0.1499 - acc: 0.9622
- val_loss: 0.0998 - val_acc: 0.9756
Epoch 7/20
60000/60000 [=====] - 3s 53us/step - loss: 0.1418 - acc: 0.9658
- val_loss: 0.1051 - val_acc: 0.9761
Epoch 8/20
60000/60000 [=====] - 3s 54us/step - loss: 0.1364 - acc: 0.9683
- val_loss: 0.1049 - val_acc: 0.9773
Epoch 9/20
60000/60000 [=====] - 3s 53us/step - loss: 0.1353 - acc: 0.9694
- val_loss: 0.0956 - val_acc: 0.9786
Epoch 10/20
60000/60000 [=====] - 3s 52us/step - loss: 0.1315 - acc: 0.9703
- val_loss: 0.1076 - val_acc: 0.9768
Epoch 11/20
60000/60000 [=====] - 3s 53us/step - loss: 0.1315 - acc: 0.9708
- val_loss: 0.1046 - val_acc: 0.9795
Epoch 12/20
60000/60000 [=====] - 3s 54us/step - loss: 0.1253 - acc: 0.9726
- val_loss: 0.1084 - val_acc: 0.9780
Epoch 13/20
60000/60000 [=====] - 3s 54us/step - loss: 0.1214 - acc: 0.9734
- val_loss: 0.1086 - val_acc: 0.9803
Epoch 14/20
60000/60000 [=====] - 3s 53us/step - loss: 0.1226 - acc: 0.9744
- val_loss: 0.1023 - val_acc: 0.9795
Epoch 15/20
60000/60000 [=====] - 3s 54us/step - loss: 0.1199 - acc: 0.9748
- val_loss: 0.1110 - val_acc: 0.9791
Epoch 16/20
60000/60000 [=====] - 3s 53us/step - loss: 0.1214 - acc: 0.9748
- val_loss: 0.1177 - val_acc: 0.9808
Epoch 17/20
60000/60000 [=====] - 3s 55us/step - loss: 0.1274 - acc: 0.9743
- val_loss: 0.1164 - val_acc: 0.9793
Epoch 18/20
60000/60000 [=====] - 3s 55us/step - loss: 0.1237 - acc: 0.9761
- val_loss: 0.1234 - val_acc: 0.9777
Epoch 19/20
60000/60000 [=====] - 3s 55us/step - loss: 0.1228 - acc: 0.9755
- val_loss: 0.1066 - val_acc: 0.9812
Epoch 20/20
60000/60000 [=====] - 3s 56us/step - loss: 0.1218 - acc: 0.9767
- val_loss: 0.1194 - val_acc: 0.9811
Epoch 20/20
60000/60000 [=====] - 3s 55us/step - loss: 0.1302 - acc: 0.9756
- val_loss: 0.1145 - val_acc: 0.9794

```

In [16]:

```
model4.summary()
```

Model: "sequential\_2"

Layer (type)	Output Shape	Param #
dense_5 (Dense)	(None, 512)	401920
dropout_4 (Dropout)	(None, 512)	0
dense_6 (Dense)	(None, 256)	131328
dropout_5 (Dropout)	(None, 256)	0
dense_7 (Dense)	(None, 128)	32896
dropout_6 (Dropout)	(None, 128)	0
dense_8 (Dense)	(None, 10)	1290
Total params: 567,434		
Trainable params: 567,434		

trainable params: 60,710  
Non-trainable params: 0

---

In [17]:

```
score = model4.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.11451425200372187

Test accuracy: 0.9794



**Good stats seen all along the epochs.**

In [18]:

```
w_after = model4.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)

fig = plt.figure(figsize=(10,7))
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w, color='b')
plt.xlabel('Hidden Layer 1')
```

```
plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='r')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w, color='y')
plt.xlabel('Output Layer ')
plt.show()
```



## 4.2 RELU ACTIVATION WITH BATCH NORM AND DROPOUT(0.3) Optimiser-(RMSprop)

In [0]:

```
from keras.layers.normalization import BatchNormalization
model4 = Sequential()

model4.add(Dense(512, input_dim=input_dim, activation='relu'))
model4.add(BatchNormalization())
model4.add(Dropout(0.3))

model4.add(Dense(256, input_dim=input_dim, activation='relu'))
model4.add(BatchNormalization())
model4.add(Dropout(0.3))

model4.add(Dense(128, input_dim=input_dim, activation='relu'))
model4.add(BatchNormalization())
model4.add(Dropout(0.3))

model4.add(Dense(10, input_dim=input_dim, activation='softmax'))
```

In [20]:

```
model4.compile(optimizer='RMSprop', loss='categorical_crossentropy', metrics=['accuracy'])

history = model4.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1)
```

```
, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 7s 116us/step - loss: 0.2904 - acc: 0.9130  
- val\_loss: 0.1020 - val\_acc: 0.9680

Epoch 2/20

60000/60000 [=====] - 6s 103us/step - loss: 0.1353 - acc: 0.9599  
- val\_loss: 0.0881 - val\_acc: 0.9724

Epoch 3/20

60000/60000 [=====] - 6s 103us/step - loss: 0.1040 - acc: 0.9687  
- val\_loss: 0.0740 - val\_acc: 0.9779

Epoch 4/20

60000/60000 [=====] - 6s 102us/step - loss: 0.0854 - acc: 0.9735  
- val\_loss: 0.0634 - val\_acc: 0.9809

Epoch 5/20

60000/60000 [=====] - 6s 102us/step - loss: 0.0745 - acc: 0.9778  
- val\_loss: 0.0631 - val\_acc: 0.9819

Epoch 6/20

60000/60000 [=====] - 6s 101us/step - loss: 0.0641 - acc: 0.9803  
- val\_loss: 0.0611 - val\_acc: 0.9814

Epoch 7/20

60000/60000 [=====] - 6s 102us/step - loss: 0.0560 - acc: 0.9827  
- val\_loss: 0.0607 - val\_acc: 0.9819

Epoch 8/20

60000/60000 [=====] - 6s 102us/step - loss: 0.0522 - acc: 0.9847  
- val\_loss: 0.0626 - val\_acc: 0.9831

Epoch 9/20

60000/60000 [=====] - 6s 102us/step - loss: 0.0480 - acc: 0.9852  
- val\_loss: 0.0581 - val\_acc: 0.9825

Epoch 10/20

60000/60000 [=====] - 6s 103us/step - loss: 0.0440 - acc: 0.9863  
- val\_loss: 0.0636 - val\_acc: 0.9818

Epoch 11/20

60000/60000 [=====] - 6s 103us/step - loss: 0.0431 - acc: 0.9868  
- val\_loss: 0.0627 - val\_acc: 0.9828

Epoch 12/20

60000/60000 [=====] - 6s 101us/step - loss: 0.0366 - acc: 0.9889  
- val\_loss: 0.0626 - val\_acc: 0.9822

Epoch 13/20

60000/60000 [=====] - 6s 102us/step - loss: 0.0367 - acc: 0.9885  
- val\_loss: 0.0637 - val\_acc: 0.9815

Epoch 14/20

60000/60000 [=====] - 6s 105us/step - loss: 0.0376 - acc: 0.9884  
- val\_loss: 0.0573 - val\_acc: 0.9852

Epoch 15/20

60000/60000 [=====] - 6s 102us/step - loss: 0.0357 - acc: 0.9887  
- val\_loss: 0.0605 - val\_acc: 0.9845

Epoch 16/20

60000/60000 [=====] - 6s 100us/step - loss: 0.0314 - acc: 0.9896  
- val\_loss: 0.0632 - val\_acc: 0.9836

Epoch 17/20

60000/60000 [=====] - 6s 102us/step - loss: 0.0327 - acc: 0.9901  
- val\_loss: 0.0603 - val\_acc: 0.9846

Epoch 18/20

60000/60000 [=====] - 6s 101us/step - loss: 0.0290 - acc: 0.9908  
- val\_loss: 0.0581 - val\_acc: 0.9838

Epoch 19/20

60000/60000 [=====] - 6s 103us/step - loss: 0.0299 - acc: 0.9905  
- val\_loss: 0.0580 - val\_acc: 0.9847

Epoch 20/20

60000/60000 [=====] - 6s 103us/step - loss: 0.0264 - acc: 0.9917  
- val\_loss: 0.0585 - val\_acc: 0.9854

In [21]:

```
model4.summary()
```

Model: "sequential\_3"

Layer (type)	Output Shape	Param #
--------------	--------------	---------

dense_9 (Dense)	(None, 512)	401920
batch_normalization_1 (Batch Normalization)	(None, 512)	2048
dropout_7 (Dropout)	(None, 512)	0
dense_10 (Dense)	(None, 256)	131328
batch_normalization_2 (Batch Normalization)	(None, 256)	1024
dropout_8 (Dropout)	(None, 256)	0
dense_11 (Dense)	(None, 128)	32896
batch_normalization_3 (Batch Normalization)	(None, 128)	512
dropout_9 (Dropout)	(None, 128)	0
dense_12 (Dense)	(None, 10)	1290
=====		
Total params: 571,018		
Trainable params: 569,226		
Non-trainable params: 1,792		

In [22]:

```
score = model4.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))


# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.05845784421725511  
Test accuracy: 0.9854





The little value of dropout made the model to overfit at 6th epoch. But the stats are good till the end.

In [23]:

```
w_after = model4.get_weights()

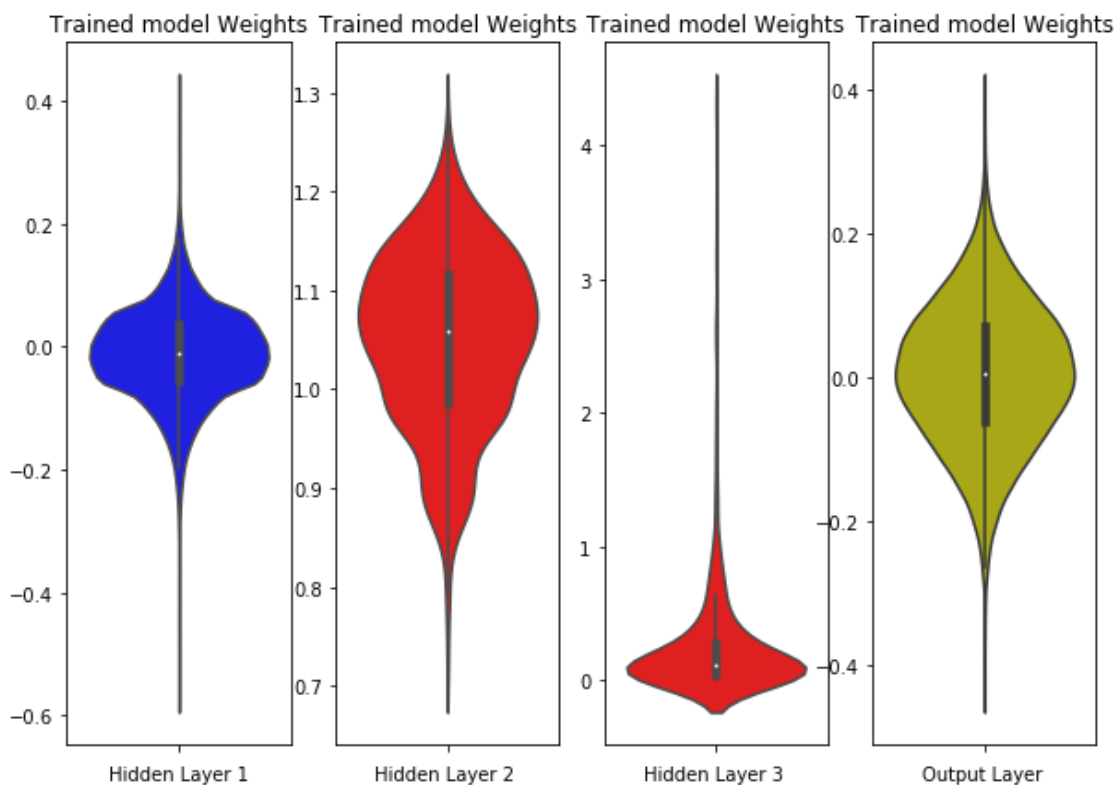
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)

fig = plt.figure(figsize=(10,7))
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='r')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



### 4.3 RELU ACTIVATION WITH BATCH NORM (Optimiser-RMSProp)

In [0]:

```
from keras.layers.normalization import BatchNormalization
model4 = Sequential()
```

```

model4.add(Dense(512, input_dim=input_dim, activation='relu'))
model4.add(BatchNormalization())

model4.add(Dense(256, input_dim=input_dim, activation='relu'))
model4.add(BatchNormalization())

model4.add(Dense(128, input_dim=input_dim, activation='relu'))
model4.add(BatchNormalization())

model4.add(Dense(10, input_dim=input_dim, activation='softmax'))

```

In [25]:

```

model4.compile(optimizer='RMSprop', loss='categorical_crossentropy', metrics=['accuracy'])

history = model4.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1,
                    validation_data=(X_test, Y_test))

```

Train on 60000 samples, validate on 10000 samples

```

Epoch 1/20
60000/60000 [=====] - 7s 111us/step - loss: 0.1710 - acc: 0.9479
- val_loss: 0.0983 - val_acc: 0.9690
Epoch 2/20
60000/60000 [=====] - 6s 95us/step - loss: 0.0748 - acc: 0.9770
- val_loss: 0.0898 - val_acc: 0.9725
Epoch 3/20
60000/60000 [=====] - 6s 95us/step - loss: 0.0495 - acc: 0.9848
- val_loss: 0.0716 - val_acc: 0.9778
Epoch 4/20
60000/60000 [=====] - 6s 95us/step - loss: 0.0369 - acc: 0.9885
- val_loss: 0.0698 - val_acc: 0.9793
Epoch 5/20
60000/60000 [=====] - 6s 94us/step - loss: 0.0270 - acc: 0.9913
- val_loss: 0.0714 - val_acc: 0.9787
Epoch 6/20
60000/60000 [=====] - 6s 94us/step - loss: 0.0239 - acc: 0.9924
- val_loss: 0.0660 - val_acc: 0.9816
Epoch 7/20
60000/60000 [=====] - 6s 96us/step - loss: 0.0178 - acc: 0.9944
- val_loss: 0.0843 - val_acc: 0.9796
Epoch 8/20
60000/60000 [=====] - 6s 96us/step - loss: 0.0167 - acc: 0.9941
- val_loss: 0.0697 - val_acc: 0.9812
Epoch 9/20
60000/60000 [=====] - 6s 96us/step - loss: 0.0134 - acc: 0.9957
- val_loss: 0.0827 - val_acc: 0.9799
Epoch 10/20
60000/60000 [=====] - 6s 95us/step - loss: 0.0121 - acc: 0.9961
- val_loss: 0.0769 - val_acc: 0.9828
Epoch 11/20
60000/60000 [=====] - 6s 97us/step - loss: 0.0114 - acc: 0.9961
- val_loss: 0.0700 - val_acc: 0.9828
Epoch 12/20
60000/60000 [=====] - 6s 95us/step - loss: 0.0104 - acc: 0.9965
- val_loss: 0.0743 - val_acc: 0.9828
Epoch 13/20
60000/60000 [=====] - 6s 98us/step - loss: 0.0096 - acc: 0.9970
- val_loss: 0.0836 - val_acc: 0.9821
Epoch 14/20
60000/60000 [=====] - 6s 96us/step - loss: 0.0080 - acc: 0.9975
- val_loss: 0.0875 - val_acc: 0.9807
Epoch 15/20
60000/60000 [=====] - 6s 97us/step - loss: 0.0065 - acc: 0.9978
- val_loss: 0.0875 - val_acc: 0.9803
Epoch 16/20
60000/60000 [=====] - 6s 95us/step - loss: 0.0078 - acc: 0.9974
- val_loss: 0.0900 - val_acc: 0.9823
Epoch 17/20
60000/60000 [=====] - 6s 96us/step - loss: 0.0065 - acc: 0.9978
- val_loss: 0.0847 - val_acc: 0.9820

```

```
Epoch 18/20
60000/60000 [=====] - 6s 94us/step - loss: 0.0078 - acc: 0.9973
- val_loss: 0.0875 - val_acc: 0.9822
Epoch 19/20
60000/60000 [=====] - 6s 96us/step - loss: 0.0064 - acc: 0.9979
- val_loss: 0.0917 - val_acc: 0.9808
Epoch 20/20
60000/60000 [=====] - 6s 94us/step - loss: 0.0068 - acc: 0.9979
- val_loss: 0.1034 - val_acc: 0.9799
```

In [26]:

```
model4.summary()
```

Model: "sequential\_4"

Layer (type)	Output Shape	Param #
dense_13 (Dense)	(None, 512)	401920
batch_normalization_4 (Batch Normalization)	(None, 512)	2048
dense_14 (Dense)	(None, 256)	131328
batch_normalization_5 (Batch Normalization)	(None, 256)	1024
dense_15 (Dense)	(None, 128)	32896
batch_normalization_6 (Batch Normalization)	(None, 128)	512
dense_16 (Dense)	(None, 10)	1290
Total params: 571,018		
Trainable params: 569,226		
Non-trainable params: 1,792		

In [27]:

```
score = model4.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.10336758222593094

Test accuracy: 0.9799





**Best case of overfitting till now.**

In [28]:

```
w_after = model4.get_weights()

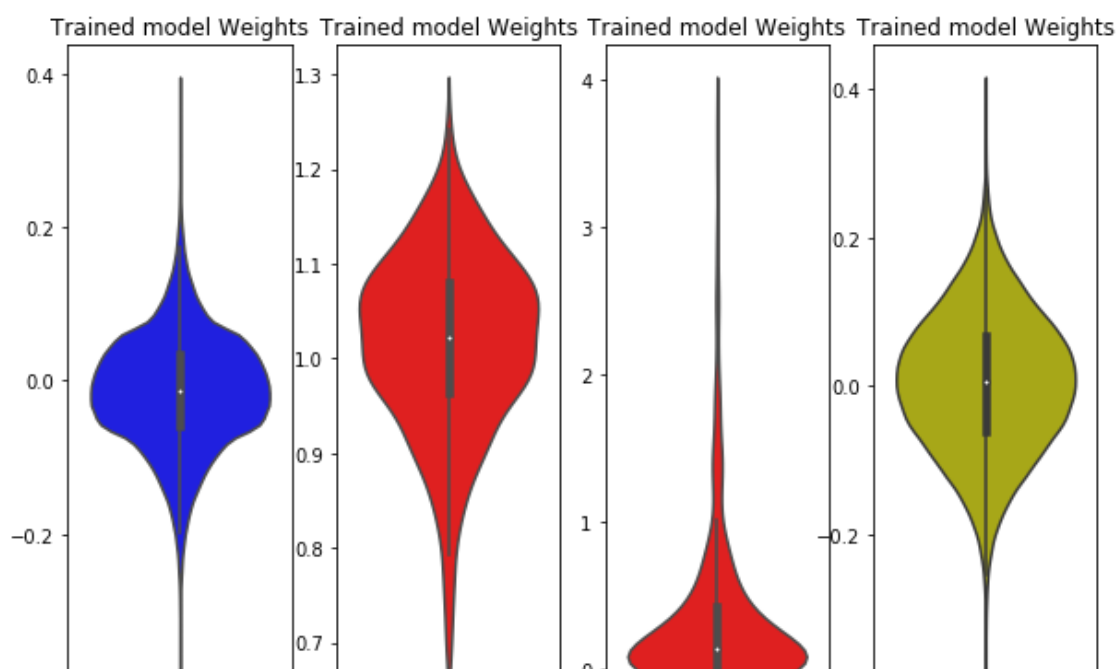
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)


fig = plt.figure(figsize=(10,7))
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='r')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```





## 5.1 SIGMOID ACTIVATION WITH DROPOUT VALUE 0.4 (Optimiser - SGD)

In [0]:

```
from keras.layers import Dropout
model5 = Sequential()
model5.add(Dense(512, input_dim=input_dim, activation='relu'))
model5.add(Dropout(0.4))
model5.add(Dense(256, input_dim=input_dim, activation='relu'))
model5.add(Dropout(0.4))
model5.add(Dense(128, input_dim=input_dim, activation='relu'))
model5.add(Dropout(0.4))
model5.add(Dense(10, input_dim=input_dim, activation='softmax'))
```

In [32]:

```
model5.compile(optimizer='SGD', loss='categorical_crossentropy', metrics=['accuracy'])

history = model5.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1,
                    validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 4s 60us/step - loss: 1.6945 - acc: 0.4247
- val_loss: 0.7545 - val_acc: 0.8054
Epoch 2/20
60000/60000 [=====] - 3s 50us/step - loss: 0.8934 - acc: 0.7086
- val_loss: 0.4545 - val_acc: 0.8809
Epoch 3/20
60000/60000 [=====] - 3s 51us/step - loss: 0.6628 - acc: 0.7939
- val_loss: 0.3635 - val_acc: 0.8979
Epoch 4/20
60000/60000 [=====] - 3s 50us/step - loss: 0.5537 - acc: 0.8311
- val_loss: 0.3132 - val_acc: 0.9098
Epoch 5/20
60000/60000 [=====] - 3s 50us/step - loss: 0.4898 - acc: 0.8530
- val_loss: 0.2830 - val_acc: 0.9185
Epoch 6/20
60000/60000 [=====] - 3s 51us/step - loss: 0.4421 - acc: 0.8677
- val_loss: 0.2596 - val_acc: 0.9258
Epoch 7/20
60000/60000 [=====] - 3s 50us/step - loss: 0.4044 - acc: 0.8799
- val_loss: 0.2398 - val_acc: 0.9323
Epoch 8/20
60000/60000 [=====] - 3s 50us/step - loss: 0.3773 - acc: 0.8892
- val_loss: 0.2242 - val_acc: 0.9358
Epoch 9/20
60000/60000 [=====] - 3s 48us/step - loss: 0.3526 - acc: 0.8969
- val_loss: 0.2114 - val_acc: 0.9388
Epoch 10/20
60000/60000 [=====] - 3s 50us/step - loss: 0.3277 - acc: 0.9035
- val_loss: 0.2009 - val_acc: 0.9419
Epoch 11/20
60000/60000 [=====] - 3s 49us/step - loss: 0.3163 - acc: 0.9077
- val_loss: 0.1889 - val_acc: 0.9438
Epoch 12/20
60000/60000 [=====] - 3s 49us/step - loss: 0.2980 - acc: 0.9135
- val_loss: 0.1803 - val_acc: 0.9468
Epoch 13/20
60000/60000 [=====] - 3s 49us/step - loss: 0.2865 - acc: 0.9159
- val_loss: 0.1740 - val_acc: 0.9486
Epoch 14/20
60000/60000 [=====] - 3s 49us/step - loss: 0.2725 - acc: 0.9204
- val_loss: 0.1666 - val_acc: 0.9505
Epoch 15/20
60000/60000 [=====] - 3s 47us/step - loss: 0.2644 - acc: 0.9234
```

```

- val_loss: 0.1596 - val_acc: 0.9514
Epoch 16/20
60000/60000 [=====] - 3s 48us/step - loss: 0.2512 - acc: 0.9270
- val_loss: 0.1553 - val_acc: 0.9532
Epoch 17/20
60000/60000 [=====] - 3s 49us/step - loss: 0.2406 - acc: 0.9304
- val_loss: 0.1481 - val_acc: 0.9554
Epoch 18/20
60000/60000 [=====] - 3s 49us/step - loss: 0.2339 - acc: 0.9313
- val_loss: 0.1443 - val_acc: 0.9553
Epoch 19/20
60000/60000 [=====] - 3s 50us/step - loss: 0.2260 - acc: 0.9342
- val_loss: 0.1392 - val_acc: 0.9568
Epoch 20/20
60000/60000 [=====] - 3s 50us/step - loss: 0.2186 - acc: 0.9363
- val_loss: 0.1357 - val_acc: 0.9590

```

In [33]:

```
model5.summary()
```

Model: "sequential\_6"

Layer (type)	Output Shape	Param #
=====		
dense_21 (Dense)	(None, 512)	401920
dropout_13 (Dropout)	(None, 512)	0
dense_22 (Dense)	(None, 256)	131328
dropout_14 (Dropout)	(None, 256)	0
dense_23 (Dense)	(None, 128)	32896
dropout_15 (Dropout)	(None, 128)	0
dense_24 (Dense)	(None, 10)	1290
=====		
Total params: 567,434		
Trainable params: 567,434		
Non-trainable params: 0		

In [34]:

```

score = model5.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs


vy = history.history['val_loss']
ty = history.history['loss']

```

```
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.13570419172216205

Test accuracy: 0.959



**Although it is a slow learner there is no under fitting nor overfitting.**

In [35]:

```
w_after = model5.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)


fig = plt.figure(figsize=(10,7))
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='r')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```





## 5.2 RELU ACTIVATION WITH BATCH NORM AND DROPOUT(0.3) Optimiser-(SGD)

In [0]:

```
from keras.layers.normalization import BatchNormalization
model5 = Sequential()

model5.add(Dense(512, input_dim=input_dim, activation='relu'))
model5.add(BatchNormalization())
model5.add(Dropout(0.3))

model5.add(Dense(256, input_dim=input_dim, activation='relu'))
model5.add(BatchNormalization())
model5.add(Dropout(0.3))

model5.add(Dense(128, input_dim=input_dim, activation='relu'))
model5.add(BatchNormalization())
model5.add(Dropout(0.3))

model5.add(Dense(10, input_dim=input_dim, activation='softmax'))
```

In [37]:

```
model5.compile(optimizer='SGD', loss='categorical_crossentropy', metrics=['accuracy'])

history = model5.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1,
                    validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 7s 110us/step - loss: 0.7637 - acc: 0.7637
- val_loss: 0.2561 - val_acc: 0.9228
Epoch 2/20
60000/60000 [=====] - 5s 88us/step - loss: 0.3820 - acc: 0.8831
- val_loss: 0.1963 - val_acc: 0.9411
Epoch 3/20
60000/60000 [=====] - 5s 87us/step - loss: 0.3091 - acc: 0.9053
- val_loss: 0.1665 - val_acc: 0.9495
Epoch 4/20
60000/60000 [=====] - 5s 88us/step - loss: 0.2655 - acc: 0.9186
- val_loss: 0.1480 - val_acc: 0.9553
Epoch 5/20
60000/60000 [=====] - 5s 88us/step - loss: 0.2400 - acc: 0.9280
- val_loss: 0.1334 - val_acc: 0.9599
Epoch 6/20
60000/60000 [=====] - 5s 89us/step - loss: 0.2181 - acc: 0.9343
- val_loss: 0.1250 - val_acc: 0.9617
Epoch 7/20
60000/60000 [=====] - 5s 89us/step - loss: 0.2011 - acc: 0.9397
- val_loss: 0.1174 - val_acc: 0.9644
Epoch 8/20
60000/60000 [=====] - 5s 88us/step - loss: 0.1863 - acc: 0.9437
- val_loss: 0.1116 - val_acc: 0.9671
Epoch 9/20
60000/60000 [=====] - 5s 89us/step - loss: 0.1742 - acc: 0.9469
- val_loss: 0.1057 - val_acc: 0.9675
Epoch 10/20
60000/60000 [=====] - 6s 92us/step - loss: 0.1681 - acc: 0.9495
- val_loss: 0.1010 - val_acc: 0.9689
Epoch 11/20
```



```

Epoch 11/20
60000/60000 [=====] - 5s 88us/step - loss: 0.1544 - acc: 0.9541
- val_loss: 0.0969 - val_acc: 0.9701
Epoch 12/20
60000/60000 [=====] - 5s 87us/step - loss: 0.1475 - acc: 0.9554
- val_loss: 0.0932 - val_acc: 0.9719
Epoch 13/20
60000/60000 [=====] - 5s 88us/step - loss: 0.1393 - acc: 0.9568
- val_loss: 0.0904 - val_acc: 0.9723
Epoch 14/20
60000/60000 [=====] - 5s 91us/step - loss: 0.1368 - acc: 0.9580
- val_loss: 0.0862 - val_acc: 0.9739
Epoch 15/20
60000/60000 [=====] - 5s 89us/step - loss: 0.1292 - acc: 0.9606
- val_loss: 0.0838 - val_acc: 0.9739
Epoch 16/20
60000/60000 [=====] - 5s 89us/step - loss: 0.1238 - acc: 0.9619
- val_loss: 0.0824 - val_acc: 0.9747
Epoch 17/20
60000/60000 [=====] - 5s 89us/step - loss: 0.1200 - acc: 0.9626
- val_loss: 0.0809 - val_acc: 0.9753
Epoch 18/20
60000/60000 [=====] - 5s 91us/step - loss: 0.1132 - acc: 0.9653
- val_loss: 0.0780 - val_acc: 0.9755
Epoch 19/20
60000/60000 [=====] - 6s 95us/step - loss: 0.1096 - acc: 0.9660
- val_loss: 0.0770 - val_acc: 0.9766
Epoch 20/20
60000/60000 [=====] - 6s 92us/step - loss: 0.1047 - acc: 0.9675
- val_loss: 0.0760 - val_acc: 0.9763

```

In [38]:

```
model5.summary()
```

Model: "sequential\_7"

Layer (type)	Output Shape	Param #
dense_25 (Dense)	(None, 512)	401920
batch_normalization_7 (Batch Normalization)	(None, 512)	2048
dropout_16 (Dropout)	(None, 512)	0
dense_26 (Dense)	(None, 256)	131328
batch_normalization_8 (Batch Normalization)	(None, 256)	1024
dropout_17 (Dropout)	(None, 256)	0
dense_27 (Dense)	(None, 128)	32896
batch_normalization_9 (Batch Normalization)	(None, 128)	512
dropout_18 (Dropout)	(None, 128)	0
dense_28 (Dense)	(None, 10)	1290
Total params: 571,018		
Trainable params: 569,226		
Non-trainable params: 1,792		

In [39]:

```

score = model5.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

```

```

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.07599864422457758

Test accuracy: 0.9763



**The little value of dropout made the model to learn faster than the previous model and all the stats are good till the end.**

In [40]:

```

w_after = model5.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)

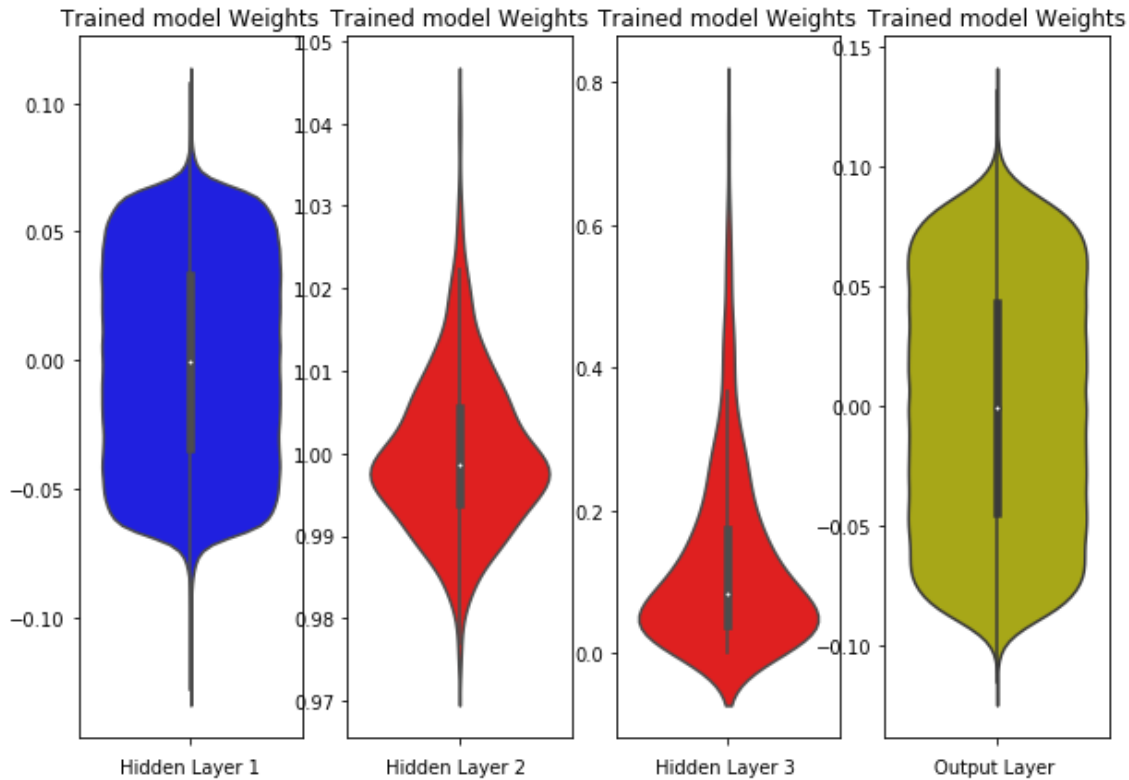
fig = plt.figure(figsize=(10,7))
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='r')
plt.xlabel('Hidden Layer 3 ')

```

```
plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w, color='y')
plt.xlabel('Output Layer ')
plt.show()
```



### 5.3 RELU ACTIVATION WITH BATCH NORM (Optimiser-SGD)

In [0]:

```
from keras.layers.normalization import BatchNormalization
model5 = Sequential()

model5.add(Dense(512, input_dim=input_dim, activation='relu'))
model5.add(BatchNormalization())

model5.add(Dense(256, input_dim=input_dim, activation='relu'))
model5.add(BatchNormalization())

model5.add(Dense(128, input_dim=input_dim, activation='relu'))
model5.add(BatchNormalization())

model5.add(Dense(10, input_dim=input_dim, activation='softmax'))
```

In [42]:

```
model5.compile(optimizer='SGD', loss='categorical_crossentropy', metrics=['accuracy'])

history = model5.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1,
                    validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 6s 108us/step - loss: 0.4255 - acc: 0.8754  
- val\_loss: 0.2161 - val\_acc: 0.9380

Epoch 2/20

60000/60000 [=====] - 5s 84us/step - loss: 0.1857 - acc: 0.9478  
- val\_loss: 0.1634 - val\_acc: 0.9521

Epoch 3/20

60000/60000 [=====] - 5s 86us/step - loss: 0.1367 - acc: 0.9627  
- val\_loss: 0.1384 - val\_acc: 0.9593

Epoch 4/20

60000/60000 [=====] - 5s 86us/step - loss: 0.1096 - acc: 0.9708  
- val\_loss: 0.1243 - val\_acc: 0.9636

```

Epoch 5/20
60000/60000 [=====] - 5s 89us/step - loss: 0.0915 - acc: 0.9758
- val_loss: 0.1153 - val_acc: 0.9649
Epoch 6/20
60000/60000 [=====] - 5s 88us/step - loss: 0.0771 - acc: 0.9805
- val_loss: 0.1084 - val_acc: 0.9663
Epoch 7/20
60000/60000 [=====] - 5s 87us/step - loss: 0.0652 - acc: 0.9842
- val_loss: 0.1027 - val_acc: 0.9681
Epoch 8/20
60000/60000 [=====] - 5s 87us/step - loss: 0.0573 - acc: 0.9860
- val_loss: 0.0985 - val_acc: 0.9687
Epoch 9/20
60000/60000 [=====] - 5s 83us/step - loss: 0.0493 - acc: 0.9886
- val_loss: 0.0954 - val_acc: 0.9698
Epoch 10/20
60000/60000 [=====] - 5s 87us/step - loss: 0.0429 - acc: 0.9906
- val_loss: 0.0923 - val_acc: 0.9705
Epoch 11/20
60000/60000 [=====] - 5s 87us/step - loss: 0.0380 - acc: 0.9919
- val_loss: 0.0906 - val_acc: 0.9713
Epoch 12/20
60000/60000 [=====] - 5s 88us/step - loss: 0.0333 - acc: 0.9932
- val_loss: 0.0893 - val_acc: 0.9715
Epoch 13/20
60000/60000 [=====] - 5s 89us/step - loss: 0.0297 - acc: 0.9945
- val_loss: 0.0873 - val_acc: 0.9723
Epoch 14/20
60000/60000 [=====] - 5s 90us/step - loss: 0.0261 - acc: 0.9956
- val_loss: 0.0865 - val_acc: 0.9725
Epoch 15/20
60000/60000 [=====] - 5s 87us/step - loss: 0.0236 - acc: 0.9961
- val_loss: 0.0846 - val_acc: 0.9714
Epoch 16/20
60000/60000 [=====] - 5s 87us/step - loss: 0.0213 - acc: 0.9974
- val_loss: 0.0844 - val_acc: 0.9723
Epoch 17/20
60000/60000 [=====] - 5s 85us/step - loss: 0.0195 - acc: 0.9975
- val_loss: 0.0841 - val_acc: 0.9723
Epoch 18/20
60000/60000 [=====] - 5s 86us/step - loss: 0.0177 - acc: 0.9979
- val_loss: 0.0829 - val_acc: 0.9726
Epoch 19/20
60000/60000 [=====] - 5s 86us/step - loss: 0.0163 - acc: 0.9983
- val_loss: 0.0830 - val_acc: 0.9735
Epoch 20/20
60000/60000 [=====] - 5s 86us/step - loss: 0.0147 - acc: 0.9986
- val_loss: 0.0824 - val_acc: 0.9741

```

In [43]:

```
model5.summary()
```

Model: "sequential\_8"

Layer (type)	Output Shape	Param #
dense_29 (Dense)	(None, 512)	401920
batch_normalization_10 (Batch Normalization)	(None, 512)	2048
dense_30 (Dense)	(None, 256)	131328
batch_normalization_11 (Batch Normalization)	(None, 256)	1024
dense_31 (Dense)	(None, 128)	32896
batch_normalization_12 (Batch Normalization)	(None, 128)	512
dense_32 (Dense)	(None, 10)	1290
Total params: 571.018		

Trainable params: 569,226  
Non-trainable params: 1,792

---

In [44]:

```
score = model5.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.08236163575346582

Test accuracy: 0.9741



**The validation loss is little slow in decreasing.**

In [45]:

```
w_after = model5.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)

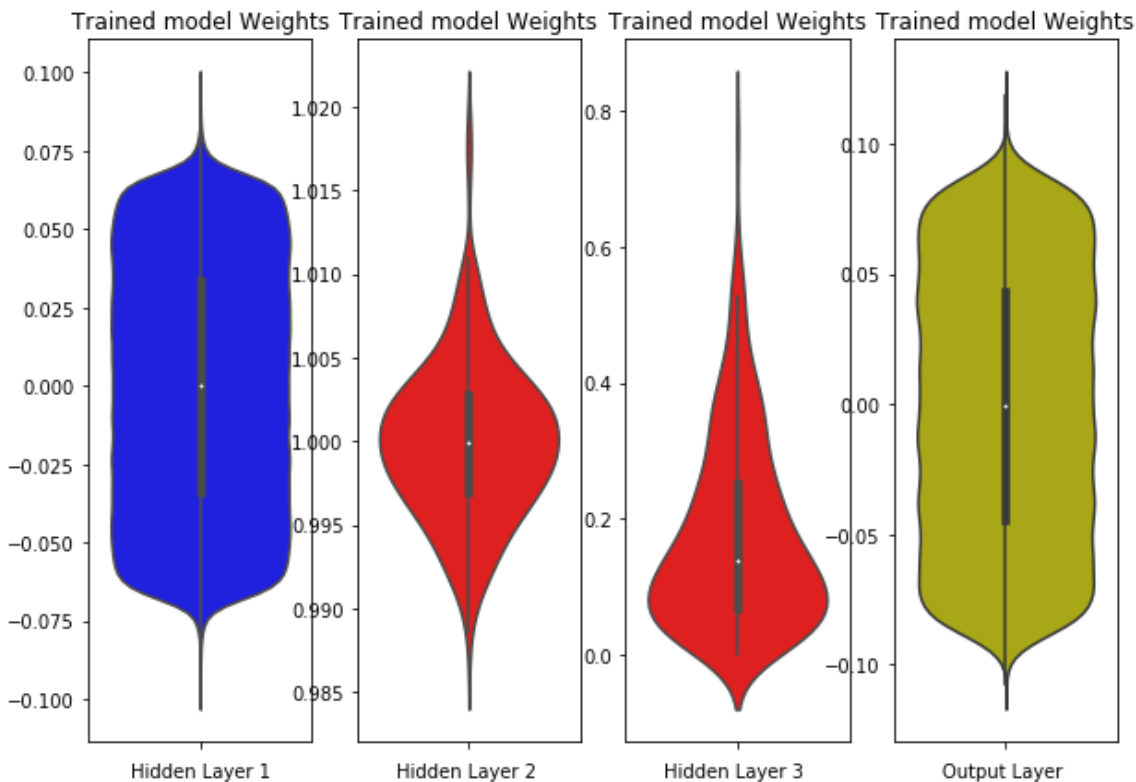
fig = plt.figure(figsize=(10,7))
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w, color='b')
```

```
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='r')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w, color='y')
plt.xlabel('Output Layer ')
plt.show()
```



## 6.1 RELU ACTIVATION WITH DROPOUT VALUE 0.5 (Optimiser - Adamax)

In [0]:

```
from keras.layers import Dropout
model5 = Sequential()
model5.add(Dense(512, input_dim=input_dim, activation='relu'))
model5.add(Dropout(0.5))
model5.add(Dense(256, input_dim=input_dim, activation='relu'))
model5.add(Dropout(0.5))
model5.add(Dense(128, input_dim=input_dim, activation='relu'))
model5.add(Dropout(0.5))
model5.add(Dense(10, input_dim=input_dim, activation='softmax'))
```

In [47]:

```
model5.compile(optimizer='Adamax', loss='categorical_crossentropy', metrics=['accuracy'])

history = model5.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1,
                    validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 4s 74us/step - loss: 0.6029 - acc: 0.8106
- val_loss: 0.1855 - val_acc: 0.9462
Epoch 2/20
```

```

Epoch 2/20
60000/60000 [=====] - 3s 57us/step - loss: 0.2661 - acc: 0.9249
- val_loss: 0.1303 - val_acc: 0.9616
Epoch 3/20
60000/60000 [=====] - 3s 56us/step - loss: 0.2040 - acc: 0.9435
- val_loss: 0.1082 - val_acc: 0.9684
Epoch 4/20
60000/60000 [=====] - 3s 58us/step - loss: 0.1697 - acc: 0.9525
- val_loss: 0.1010 - val_acc: 0.9697
Epoch 5/20
60000/60000 [=====] - 3s 56us/step - loss: 0.1498 - acc: 0.9588
- val_loss: 0.0895 - val_acc: 0.9731
Epoch 6/20
60000/60000 [=====] - 3s 57us/step - loss: 0.1317 - acc: 0.9628
- val_loss: 0.0844 - val_acc: 0.9750
Epoch 7/20
60000/60000 [=====] - 3s 57us/step - loss: 0.1191 - acc: 0.9661
- val_loss: 0.0801 - val_acc: 0.9757
Epoch 8/20
60000/60000 [=====] - 4s 58us/step - loss: 0.1090 - acc: 0.9690
- val_loss: 0.0767 - val_acc: 0.9775
Epoch 9/20
60000/60000 [=====] - 3s 56us/step - loss: 0.1007 - acc: 0.9721
- val_loss: 0.0751 - val_acc: 0.9790
Epoch 10/20
60000/60000 [=====] - 3s 57us/step - loss: 0.0942 - acc: 0.9738
- val_loss: 0.0755 - val_acc: 0.9792
Epoch 11/20
60000/60000 [=====] - 3s 58us/step - loss: 0.0877 - acc: 0.9749
- val_loss: 0.0742 - val_acc: 0.9793
Epoch 12/20
60000/60000 [=====] - 4s 59us/step - loss: 0.0818 - acc: 0.9773
- val_loss: 0.0689 - val_acc: 0.9814
Epoch 13/20
60000/60000 [=====] - 4s 59us/step - loss: 0.0779 - acc: 0.9777
- val_loss: 0.0698 - val_acc: 0.9809
Epoch 14/20
60000/60000 [=====] - 3s 57us/step - loss: 0.0733 - acc: 0.9784
- val_loss: 0.0678 - val_acc: 0.9814
Epoch 15/20
60000/60000 [=====] - 3s 57us/step - loss: 0.0709 - acc: 0.9797
- val_loss: 0.0687 - val_acc: 0.9816
Epoch 16/20
60000/60000 [=====] - 3s 56us/step - loss: 0.0670 - acc: 0.9805
- val_loss: 0.0688 - val_acc: 0.9813
Epoch 17/20
60000/60000 [=====] - 3s 57us/step - loss: 0.0648 - acc: 0.9810
- val_loss: 0.0694 - val_acc: 0.9808
Epoch 18/20
60000/60000 [=====] - 3s 57us/step - loss: 0.0612 - acc: 0.9817
- val_loss: 0.0661 - val_acc: 0.9830
Epoch 19/20
60000/60000 [=====] - 3s 57us/step - loss: 0.0582 - acc: 0.9824
- val_loss: 0.0668 - val_acc: 0.9823
Epoch 20/20
60000/60000 [=====] - 3s 57us/step - loss: 0.0561 - acc: 0.9841
- val_loss: 0.0660 - val_acc: 0.9832

```

In [48]:

```
model5.summary()
```

Model: "sequential\_9"

Layer (type)	Output Shape	Param #
dense_33 (Dense)	(None, 512)	401920
dropout_19 (Dropout)	(None, 512)	0
dense_34 (Dense)	(None, 256)	131328
dropout_20 (Dropout)	(None, 256)	0

dense_35 (Dense)	(None, 128)	32896
dropout_21 (Dropout)	(None, 128)	0
dense_36 (Dense)	(None, 10)	1290
=====		
Total params: 567,434		
Trainable params: 567,434		
Non-trainable params: 0		

In [49]:

```
score = model5.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))


# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.06598338727326336  
Test accuracy: 0.9832



Good stats seen all along the epochs.Overfitting at 15th epoch.

In [50]:

```
w_after = model5.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
```



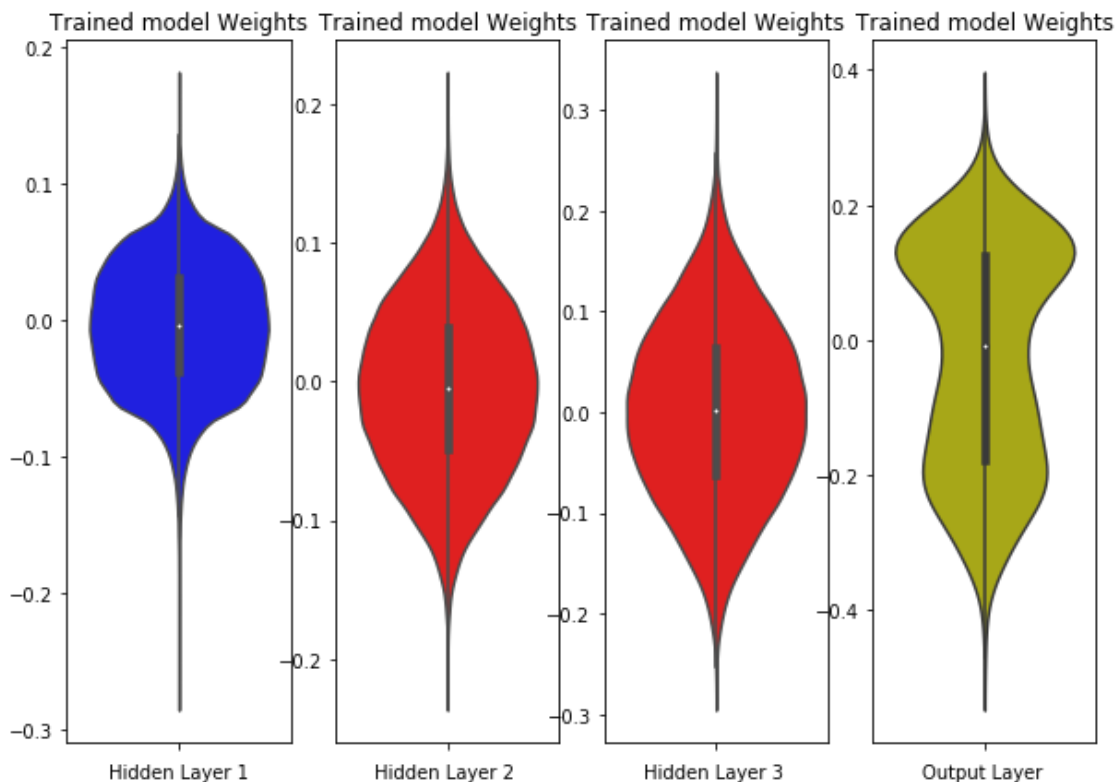
```
out_w = w_after[6].flatten().reshape(-1,1)
```

```
fig = plt.figure(figsize=(10,7))
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='r')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



## 6.2 RELU ACTIVATION WITH BATCH NORM AND DROPOUT(0.3) Optimiser-(Adamax)

In [0]:

```
from keras.layers.normalization import BatchNormalization
model6 = Sequential()

model6.add(Dense(512, input_dim=input_dim, activation='relu'))
model6.add(BatchNormalization())
model6.add(Dropout(0.3))

model6.add(Dense(256, input_dim=input_dim, activation='relu'))
model6.add(BatchNormalization())
model6.add(Dropout(0.3))

model6.add(Dense(128, input_dim=input_dim, activation='relu'))
model6.add(BatchNormalization())
model6.add(Dropout(0.3))
```

```
model6.add(Dense(10, input_dim=input_dim, activation='softmax'))
```

In [52]:

```
model6.compile(optimizer='Adamax', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model6.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 8s 130us/step - loss: 0.3441 - acc: 0.8962  
- val\_loss: 0.1332 - val\_acc: 0.9567

Epoch 2/20

60000/60000 [=====] - 6s 99us/step - loss: 0.1713 - acc: 0.9481  
- val\_loss: 0.0953 - val\_acc: 0.9702

Epoch 3/20

60000/60000 [=====] - 6s 100us/step - loss: 0.1265 - acc: 0.9623  
- val\_loss: 0.0868 - val\_acc: 0.9726

Epoch 4/20

60000/60000 [=====] - 6s 100us/step - loss: 0.1040 - acc: 0.9677  
- val\_loss: 0.0798 - val\_acc: 0.9739

Epoch 5/20

60000/60000 [=====] - 6s 100us/step - loss: 0.0876 - acc: 0.9735  
- val\_loss: 0.0689 - val\_acc: 0.9787

Epoch 6/20

60000/60000 [=====] - 6s 99us/step - loss: 0.0743 - acc: 0.9771  
- val\_loss: 0.0655 - val\_acc: 0.9795

Epoch 7/20

60000/60000 [=====] - 6s 100us/step - loss: 0.0653 - acc: 0.9792  
- val\_loss: 0.0637 - val\_acc: 0.9815

Epoch 8/20

60000/60000 [=====] - 6s 100us/step - loss: 0.0590 - acc: 0.9811  
- val\_loss: 0.0681 - val\_acc: 0.9799

Epoch 9/20

60000/60000 [=====] - 6s 101us/step - loss: 0.0509 - acc: 0.9836  
- val\_loss: 0.0589 - val\_acc: 0.9831

Epoch 10/20

60000/60000 [=====] - 6s 100us/step - loss: 0.0469 - acc: 0.9844  
- val\_loss: 0.0624 - val\_acc: 0.9826

Epoch 11/20

60000/60000 [=====] - 6s 99us/step - loss: 0.0410 - acc: 0.9868  
- val\_loss: 0.0606 - val\_acc: 0.9814

Epoch 12/20

60000/60000 [=====] - 6s 100us/step - loss: 0.0391 - acc: 0.9874  
- val\_loss: 0.0587 - val\_acc: 0.9825

Epoch 13/20

60000/60000 [=====] - 6s 101us/step - loss: 0.0366 - acc: 0.9878  
- val\_loss: 0.0597 - val\_acc: 0.9827

Epoch 14/20

60000/60000 [=====] - 6s 103us/step - loss: 0.0328 - acc: 0.9892  
- val\_loss: 0.0630 - val\_acc: 0.9825

Epoch 15/20

60000/60000 [=====] - 6s 101us/step - loss: 0.0327 - acc: 0.9895  
- val\_loss: 0.0623 - val\_acc: 0.9825

Epoch 16/20

60000/60000 [=====] - 6s 100us/step - loss: 0.0294 - acc: 0.9899  
- val\_loss: 0.0614 - val\_acc: 0.9839

Epoch 17/20

60000/60000 [=====] - 6s 102us/step - loss: 0.0277 - acc: 0.9908  
- val\_loss: 0.0606 - val\_acc: 0.9833

Epoch 18/20

60000/60000 [=====] - 6s 106us/step - loss: 0.0237 - acc: 0.9919  
- val\_loss: 0.0602 - val\_acc: 0.9840

Epoch 19/20

60000/60000 [=====] - 6s 106us/step - loss: 0.0269 - acc: 0.9911  
- val\_loss: 0.0591 - val\_acc: 0.9840

Epoch 20/20

60000/60000 [=====] - 6s 101us/step - loss: 0.0220 - acc: 0.9925  
- val\_loss: 0.0598 - val\_acc: 0.9839

In [53]:

```
model6.summary()
```

Model: "sequential\_10"

Layer (type)	Output Shape	Param #
dense_37 (Dense)	(None, 512)	401920
batch_normalization_13 (Batch Normalization)	(None, 512)	2048
dropout_22 (Dropout)	(None, 512)	0
dense_38 (Dense)	(None, 256)	131328
batch_normalization_14 (Batch Normalization)	(None, 256)	1024
dropout_23 (Dropout)	(None, 256)	0
dense_39 (Dense)	(None, 128)	32896
batch_normalization_15 (Batch Normalization)	(None, 128)	512
dropout_24 (Dropout)	(None, 128)	0
dense_40 (Dense)	(None, 10)	1290
Total params: 571,018		
Trainable params: 569,226		
Non-trainable params: 1,792		

In [54]:

```
score = model6.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the parameter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.059768225400491794

Test accuracy: 0.9839





The little value of dropout made the model to overfit at 7th epoch. But the stats are good till the end.

In [55]:

```
w_after = model6.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)

fig = plt.figure(figsize=(10,7))
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='r')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



### 6.3 RELU ACTIVATION WITH BATCH NORM (Optimiser-Adamax)

In [0]:

```
from keras.layers.normalization import BatchNormalization
model6 = Sequential()

model6.add(Dense(512, input_dim=input_dim, activation='relu'))
model6.add(BatchNormalization())

model6.add(Dense(256, input_dim=input_dim, activation='relu'))
model6.add(BatchNormalization())

model6.add(Dense(128, input_dim=input_dim, activation='relu'))
model6.add(BatchNormalization())

model6.add(Dense(10, input_dim=input_dim, activation='softmax'))
```

In [57]:

```
model6.compile(optimizer='Adamax', loss='categorical_crossentropy', metrics=['accuracy'])

history = model6.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1,
                    validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 8s 128us/step - loss: 0.1790 - acc: 0.9464  
- val\_loss: 0.0975 - val\_acc: 0.9699

Epoch 2/20

60000/60000 [=====] - 6s 100us/step - loss: 0.0691 - acc: 0.9797  
- val\_loss: 0.0835 - val\_acc: 0.9722

Epoch 3/20

60000/60000 [=====] - 6s 98us/step - loss: 0.0420 - acc: 0.9873  
- val\_loss: 0.0815 - val\_acc: 0.9735

Epoch 4/20

60000/60000 [=====] - 6s 97us/step - loss: 0.0276 - acc: 0.9918  
- val\_loss: 0.0720 - val\_acc: 0.9775

Epoch 5/20

60000/60000 [=====] - 6s 97us/step - loss: 0.0171 - acc: 0.9949  
- val\_loss: 0.0656 - val\_acc: 0.9791

Epoch 6/20

60000/60000 [=====] - 6s 99us/step - loss: 0.0122 - acc: 0.9964  
- val\_loss: 0.0793 - val\_acc: 0.9768

Epoch 7/20

60000/60000 [=====] - 6s 98us/step - loss: 0.0109 - acc: 0.9966  
- val\_loss: 0.0704 - val\_acc: 0.9812

Epoch 8/20

60000/60000 [=====] - 6s 97us/step - loss: 0.0074 - acc: 0.9979  
- val\_loss: 0.0691 - val\_acc: 0.9801

Epoch 9/20

60000/60000 [=====] - 6s 97us/step - loss: 0.0062 - acc: 0.9982  
- val\_loss: 0.0670 - val\_acc: 0.9809

Epoch 10/20

60000/60000 [=====] - 6s 99us/step - loss: 0.0042 - acc: 0.9989  
- val\_loss: 0.0686 - val\_acc: 0.9804

Epoch 11/20

60000/60000 [=====] - 6s 99us/step - loss: 0.0040 - acc: 0.9991  
- val\_loss: 0.0680 - val\_acc: 0.9820

Epoch 12/20

60000/60000 [=====] - 6s 98us/step - loss: 0.0036 - acc: 0.9991  
- val\_loss: 0.0753 - val\_acc: 0.9795

Epoch 13/20

60000/60000 [=====] - 6s 97us/step - loss: 0.0030 - acc: 0.9993  
- val\_loss: 0.0701 - val\_acc: 0.9806

Epoch 14/20

60000/60000 [=====] - 6s 98us/step - loss: 0.0031 - acc: 0.9991  
- val\_loss: 0.0673 - val\_acc: 0.9820

```

Epoch 15/20
60000/60000 [=====] - 6s 99us/step - loss: 0.0035 - acc: 0.9990
- val_loss: 0.0732 - val_acc: 0.9816
Epoch 16/20
60000/60000 [=====] - 6s 98us/step - loss: 0.0032 - acc: 0.9990
- val_loss: 0.0766 - val_acc: 0.9803
Epoch 17/20
60000/60000 [=====] - 6s 98us/step - loss: 0.0016 - acc: 0.9997
- val_loss: 0.0681 - val_acc: 0.9833
Epoch 18/20
60000/60000 [=====] - 6s 100us/step - loss: 0.0021 - acc: 0.9995
- val_loss: 0.0702 - val_acc: 0.9822
Epoch 19/20
60000/60000 [=====] - 6s 100us/step - loss: 0.0026 - acc: 0.9992
- val_loss: 0.0727 - val_acc: 0.9820
Epoch 20/20
60000/60000 [=====] - 6s 97us/step - loss: 0.0023 - acc: 0.9994
- val_loss: 0.0755 - val_acc: 0.9828

```

In [58]:

```
model6.summary()
```

Model: "sequential\_11"

Layer (type)	Output Shape	Param #
dense_41 (Dense)	(None, 512)	401920
batch_normalization_16 (Batch Normalization)	(None, 512)	2048
dense_42 (Dense)	(None, 256)	131328
batch_normalization_17 (Batch Normalization)	(None, 256)	1024
dense_43 (Dense)	(None, 128)	32896
batch_normalization_18 (Batch Normalization)	(None, 128)	512
dense_44 (Dense)	(None, 10)	1290
Total params: 571,018		
Trainable params: 569,226		
Non-trainable params: 1,792		

In [59]:

```

score = model6.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

```

```
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.0755472852376839

Test accuracy: 0.9828



**Overfitting has seen in the earlier epochs.**

In [60]:

```
w_after = model6.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
h3_w = w_after[4].flatten().reshape(-1,1)
out_w = w_after[6].flatten().reshape(-1,1)


fig = plt.figure(figsize=(10,7))
plt.title("Weight matrices after model trained")
plt.subplot(1, 4, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 4, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 4, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h3_w, color='r')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```





In [0]:

```
from prettytable import PrettyTable
F='Not present'
T='Present'
Model = ["2","2","2","2","2","3","3","3","3","5","5","5","5"]
test_scr=[0.08,0.09,0.06,0.08,0.06,0.108,0.06,0.07,0.05,0.08,0.07,0.10,0.06]
test_acc=[0.98,0.98,0.98,0.97,0.98,0.97,0.98,0.98,0.98,0.98,0.98,0.97,0.98]
d=[F,F,T,F,T,F,T,F,T,F,T,F,T]
b=[F,F,F,T,T,F,F,T,T,F,F,T,T]
sno =[1,2,3,4,5,6,7,8,9,10,11,12,13]
table = PrettyTable()
table.add_column('S-NO',sno)
table.add_column("Layers",Model)
table.add_column("Test Loss",test_scr)
table.add_column("Test Acuracy",test_acc)
table.add_column("Dropout Present",d)
table.add_column("Batch Norm Present",b)
```

In [0]:

```
print(table)
print('Dropout is 0.5 in all models')
```

S-NO	Layers	Test Loss	Test Acuracy	Dropout Present	Batch Norm Present
1	2	0.08	0.98	Not present	Not present
2	2	0.09	0.98	Not present	Not present
3	2	0.06	0.98	Present	Not present
4	2	0.08	0.97	Not present	Present
5	2	0.06	0.98	Present	Present
6	3	0.108	0.97	Not present	Not present
7	3	0.06	0.98	Present	Not present
8	3	0.07	0.98	Not present	Present
9	3	0.05	0.98	Present	Present
10	5	0.08	0.98	Not present	Not present
11	5	0.07	0.98	Present	Not present
12	5	0.1	0.97	Not present	Present
13	5	0.06	0.98	Present	Present

Dropout is 0.5 in all models

In [0]:

```
from prettytable import PrettyTable
F='Not present'
T='Present'
T5='Present (Value=0.5)'
T3='Present (Value=0.3)'
T4='Present (Value=0.4)'
Optimiser = ["RMSprop","RMSprop","RMSprop","SGD","SGD","SGD","Adamax","Adamax","Adamax"]
test_scr=[0.114,0.05,0.103,0.135,0.07,0.082,0.065,0.059,0.075]
test_acc=[0.979,0.985,0.979,0.959,0.976,0.974,0.983,0.983,0.982]
d=[T5,T3,F,T4,T3,F,T5,T3,F]
b=[F,T,T,F,T,T,F,T,T]
sno =[1,2,3,4,5,6,7,8,9]
```



```

table = PrettyTable()
table.add_column('S-NO', sno)
table.add_column("Optimizer", Optimiser)
table.add_column("Test Loss", test_scr)
table.add_column("Test Acuracy", test_acc)
table.add_column("Dropout Present", d)
table.add_column("Batch Norm Present", b)

```

In [62]:

```

print('Pretty table of models with different optimizers other than ADAM')
print(table)

```

Pretty table of models with different optimizers other than ADAM

```

+-----+-----+-----+-----+-----+-----+
--+
| S-NO | Optimizer | Test Loss | Test Acuracy | Dropout Present | Batch Norm Present |
|
+-----+-----+-----+-----+-----+-----+
--+
| 1 | RMSprop | 0.114 | 0.979 | Present (Value=0.5) | Not present |
|
| 2 | RMSprop | 0.05 | 0.985 | Present (Value=0.3) | Present |
|
| 3 | RMSprop | 0.103 | 0.979 | Not present | Present |
|
| 4 | SGD | 0.135 | 0.959 | Present (Value=0.4) | Not present |
|
| 5 | SGD | 0.07 | 0.976 | Present (Value=0.3) | Present |
|
| 6 | SGD | 0.082 | 0.974 | Not present | Present |
|
| 7 | Adamax | 0.065 | 0.983 | Present (Value=0.5) | Not present |
|
| 8 | Adamax | 0.059 | 0.983 | Present (Value=0.3) | Present |
|
| 9 | Adamax | 0.075 | 0.982 | Not present | Present |
|
+-----+-----+-----+-----+-----+-----+
--+

```

**Models with RMSprop and Adamax as their optimizers gave better Test Accuracy.**

## Conclusions

- 1. Naive models (without dropout or batch norm) failed as they ended up in overfitting and higher validation loss.**
- 2. Dropout has clearly played a vital role in minimising the validation loss to as low as training loss after the last epoch.**
- 3. Models with Batch normalisation with no dropout failed in having a lower validation loss.**
- 4. Dropout with batch norm models were also good.**
- 5. Models with RMSprop optimiser gave splendid results than all other models.**
- 6. SGD optimiser models started with low accuracy and higher loss values but in the end they delivered better results. If they ran for 25 - 30 epochs ,they may achieve similar stats of best models mentioned here.**